

Console I/O and Compilation

CSC 230 : C and Software Tools

NC State Department of Computer
Science

Topics for Today

- Console I/O
 - Character I/O
 - Meet printf() and scanf()
- Program Execution in Java and C
- Meet the Preprocessor
- Tokenization
- Coding Style

Console I/O In C

- In C, I/O is provided by functions in the *standard library*
 - This library is expected on all platforms
- To use the I/O parts of the standard library, you need to include the header file:

```
#include <stdio.h>
```

- We'll also get some use out of:

```
#include <stdlib.h>
```

- These are *preprocessor directives*, telling the preprocessor to get these files and compile them along with our source code.

Streams

- A *stream* is a file or device we can read or write
- Just like in Java, a C program starts with three streams it can use
 - *Standard input* (input typed at the terminal)
 - *Standard output* (output to the terminal)
 - *Standard error* (more output to the terminal)
- To a program, reading and writing to the terminal looks just like reading or writing a file
 - We can even signal the End-Of-File condition on standard input:
 - Type CTRL-D if your program is running in Linux
 - Or, CTRL-Z if it's running in Windows.

Redirecting Standard Streams

- We can redirect these streams to or from actual files (without the program even noticing)
 - We won't learn about file I/O for a while, but this will let us get by without it.
- From the terminal, you can redirect standard input from a file

```
$ myProgram < input_file.txt
```

- ... or standard output to a file.

```
$ myProgram > output_file.txt
```

Reading just one Character

- stdio.h provides a function
`int getchar(void)`

It returns an int, the next character read.

It doesn't take any parameters.

- Returns the next character read (as a small non-negative integer)
- ... or, if there's no more input, the constant EOF (the value -1)
- That's why its return type is int instead of char (students seem to forget this)

Remember

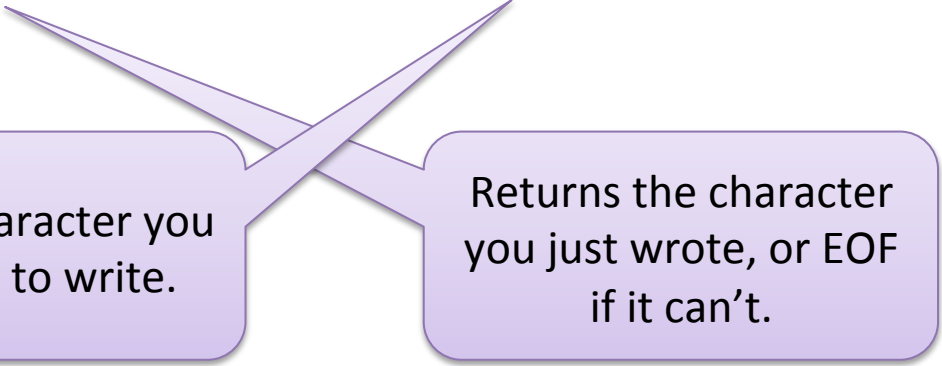
getchar()

returns int

Writing just one Character

- `stdio.h` also provides

```
int putchar( int c )
```



The character you
want to write.

Returns the character
you just wrote, or EOF
if it can't.

Character I/O Example

```
#include <stdio.h>
#include <stdlib.h>
```

```
int main()
{
    int ch = getchar();
    while ( ch != EOF ) {
        if ( ch == ' ' )
            putchar( '-' );
        else
            putchar( ch );
        ch = getchar();
    }
}
```

```
return EXIT_SUCCESS;
}
```

Read just one character.

See if we succeeded.

Print just one character.

A named constant for success!

Formatted Output

- The printf() function is good for generating formatted output
- You probably saw a similar function in Java. It works like:

```
printf( "value: %6.2f\n", 3.1415926 );
```

This is a *format string*.
Most of it gets printed
literally.

... but not parts like
this. This is a
conversion specification.

Each conversion
specification says how to
print one of the
remaining parameters.

Making Sense of Conversion Specifications

This says “Here comes a conversion specification”

This last character (or characters) says what to print, here it’s a double.

`%6.2f`

Minimum *Field Width*
This part is optional. It says “use at least 6 characters of output”.

Precision
This part is optional. For real numbers, it says “round output two fractional digits”.

Conversion Specification Examples

- There are lots of ways to print a value like 33.3:

Conversion Specification	Output
%7.1f	33.3
%14.10f	33.3000000000
%20.20f	33.299999999999999715783

3 spaces

1 space

no spaces

Conversion Specification Overview

- `printf()` can output lots of different types of values.
- Here are some we'll need for now.

Format Specification	Output
<code>%c</code>	A single character (given its numeric character code)
<code>%d</code>	A decimal integer
<code>%ld</code>	A long integer
<code>%f</code>	A float or double
<code>%s</code>	A string
<code>%zd</code>	The size of some memory region (a value of type <code>size_t</code>)

Reading Formatted Input

- **printf()** is for creating formatted output to the terminal.
- Its friend, **scanf()**, is for reading formatted input.
- The **scanf()** function converts sequences of characters into other types
 - Like int, double, string, ...
 - It can skip whitespace and other input you tell it to ignore
- It's a lot like **printf()**

```
int scanf( const char *fmt, ... )
```

- It takes a format string, followed by a variable number of additional arguments
- It matches characters from the input and assigns them to the subsequent arguments.

Using scanf()

- The format string can contain one or more *conversion specifications*.
- Each of these fills in the value of one of the remaining arguments.
- You can't do this with *pass-by-value*

```
scanf( "%d", val );
```

This couldn't possibly work

- So, C uses a mechanism for *pass-by-reference*

```
scanf( "%d", &val );
```

With the &, we're passing val so scanf can change what's in it.

Using Scanf

- scanf() can match multiple conversion specifications.

```
scanf( "%f%d", &floatVal, &intVal );
```

- It returns the number of specifications matched

```
int matches = scanf( "%f%d", &floatVal, &intVal );
```

- Fewer on bad input
- ... or EOF if it reaches the end-of-file before matching any.
- And, it only changes a parameter if it can match it successfully.

Using Scanf

- So, you could do something like:

```
printf( "Please enter an integer: " );

int val;
int matches = scanf( "%d", &val );
if ( matches == 1 ) {
    printf( "You entered %d\n", val );
} else if ( matches == 0 ) {
    printf( "That's not an integer.\n" );
} else {
    printf( "You didn't enter anything.\n" );
}
```

Conversion Specification Zoo

- scanf() understands almost exactly the same conversion specifications as printf()
- ... but, there are some differences 😊

Parse input as	Conversion Specification
int	%d (in decimal)
float	%f
double	%lf
long int	%ld
A single character	%c
A whole string	%s

Matches the longest sequence of characters that looks like an int.

... longest sequence of characters that looks like a real number.

This one is a little different.

Most of these skip leading whitespace ... but not this one.

Fun with Scanf

- Consider the code:

```
float fval = 1.0f;  
int ival = 2;  
int matches = scanf( "%f%d", &fval, &ival );
```

- On the following input, what value will these three variables have?

3 4.12

- How about this input?

37.22b5

Formatted I/O Example

```
#include <stdio.h>
#include <stdlib.h>

int main()
{
    int val;

    while ( scanf( "%d", &val ) == 1 ) {
        printf( "%6d\n", val );
    }

    return EXIT_SUCCESS;
}
```

While we keep matching
one decimal integer.

Print it in a 6-character
field.

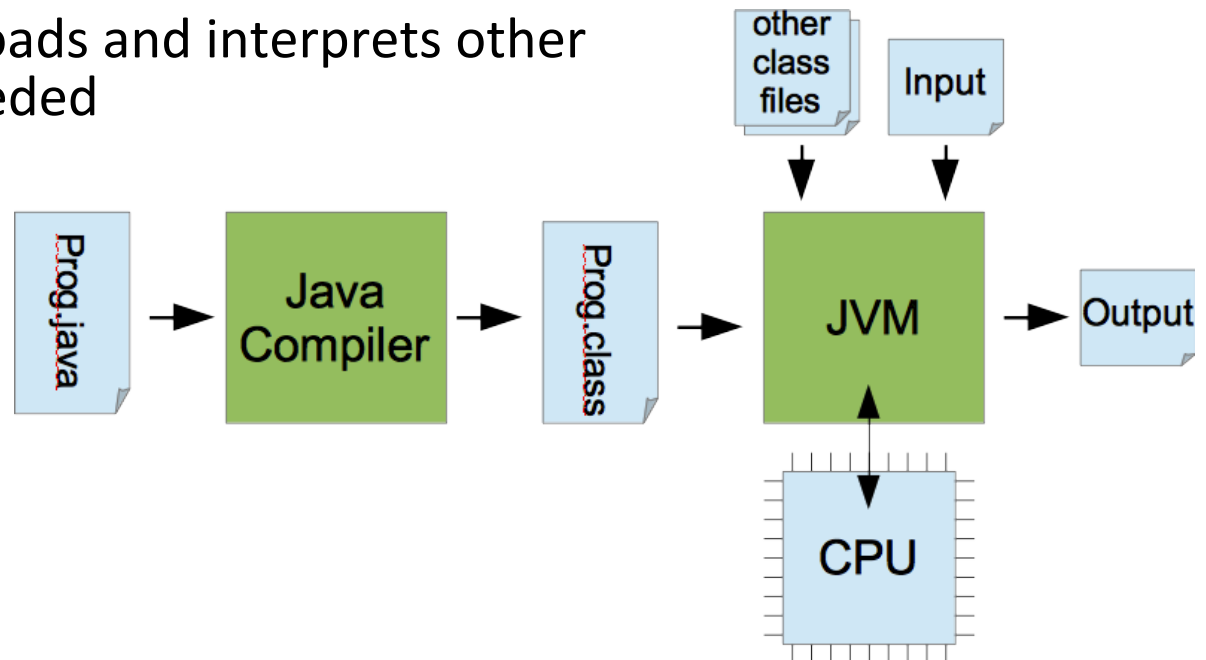
```
25
-12  77
green
62
```

output

```
25
-12
77
```

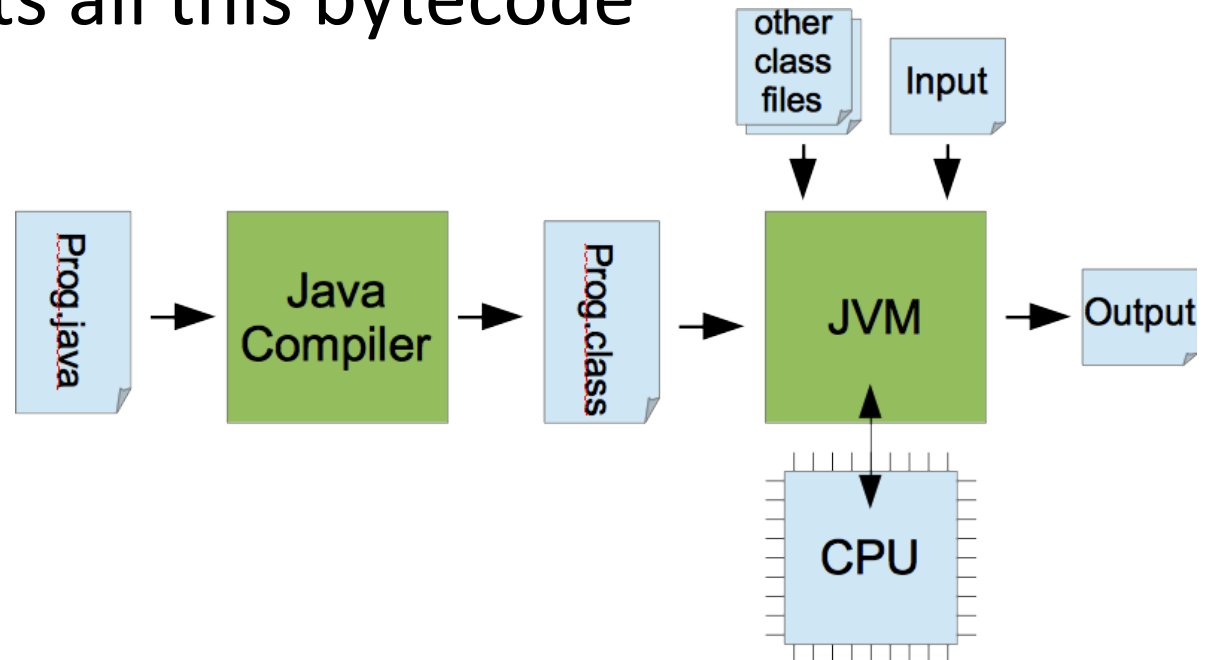
Executing Java Programs

- Java source code is *compiled* into Java class file containing *bytecode*
 - A platform-independent, intermediate representation
- To run it, we need an interpreter, the Java Virtual Machine
 - Takes a class file as input, runs native machine code to simulate each bytecode instruction
 - Automatically loads and interprets other class files as needed



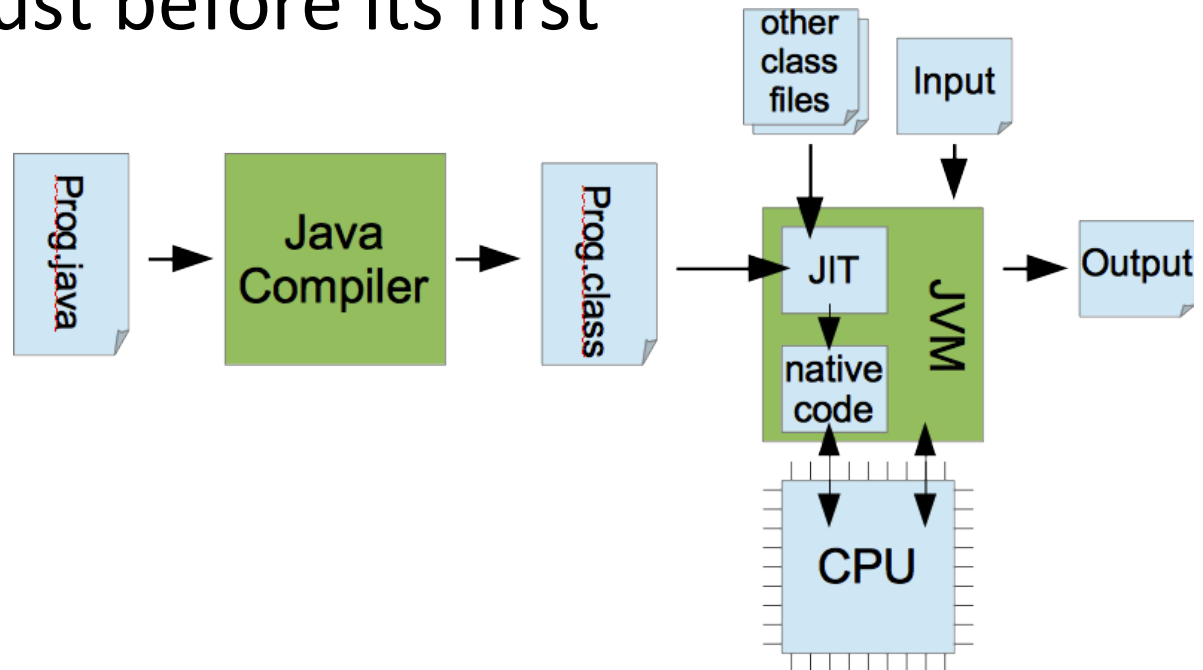
Executing Java Programs

- This is great. The class files for our compiled program are platform independent.
- But, some extra overhead is incurred as the JVM interprets all this bytecode




Executing Java Programs

- With Just-In-Time compilation, Java can get closer to native processor speeds
- Each method is compiled to native machine instructions just before its first execution




Executing C Programs

- A C Compiler generates native machine code for the target processor
- One *compilation unit* generates one *object file*




Executing C Programs

- A linker combines object files to create an executable program
 - Maybe some other objects we wrote
 - Maybe some from libraries



Executing C Programs

- The executable is ready to run
- Just load it into memory and start running at the top of main()



Compiled vs Interpreted

- Each approach has advantages.
- What do you think? Would compiled code (compiled to the native instruction set) execute faster than interpreted?
- Which would offer better support for debugging and error messages?
- Which would offer greater platform independence?
- Which would offer more opportunities for code analysis and optimization?

The Preprocessor

- The preprocessor operates on the source code before the compiler even sees it.
- Performs basic text operations
 - *Includes headers* : inserting code that enables use of code from other components
 - *Expands macros* : replacing macro names with corresponding definitions
 - More things we'll learn about later
- Lines starting with # are *preprocessor directives*
 - instructions processed (and removed) by the preprocessor

Preprocessor Constants

- Preprocessor macros give us a way to define named constants:

```
#define SIZE 25
```

Replace occurrences
of this ...

... with this.

Be careful, you probably don't
want a semi-colon here.

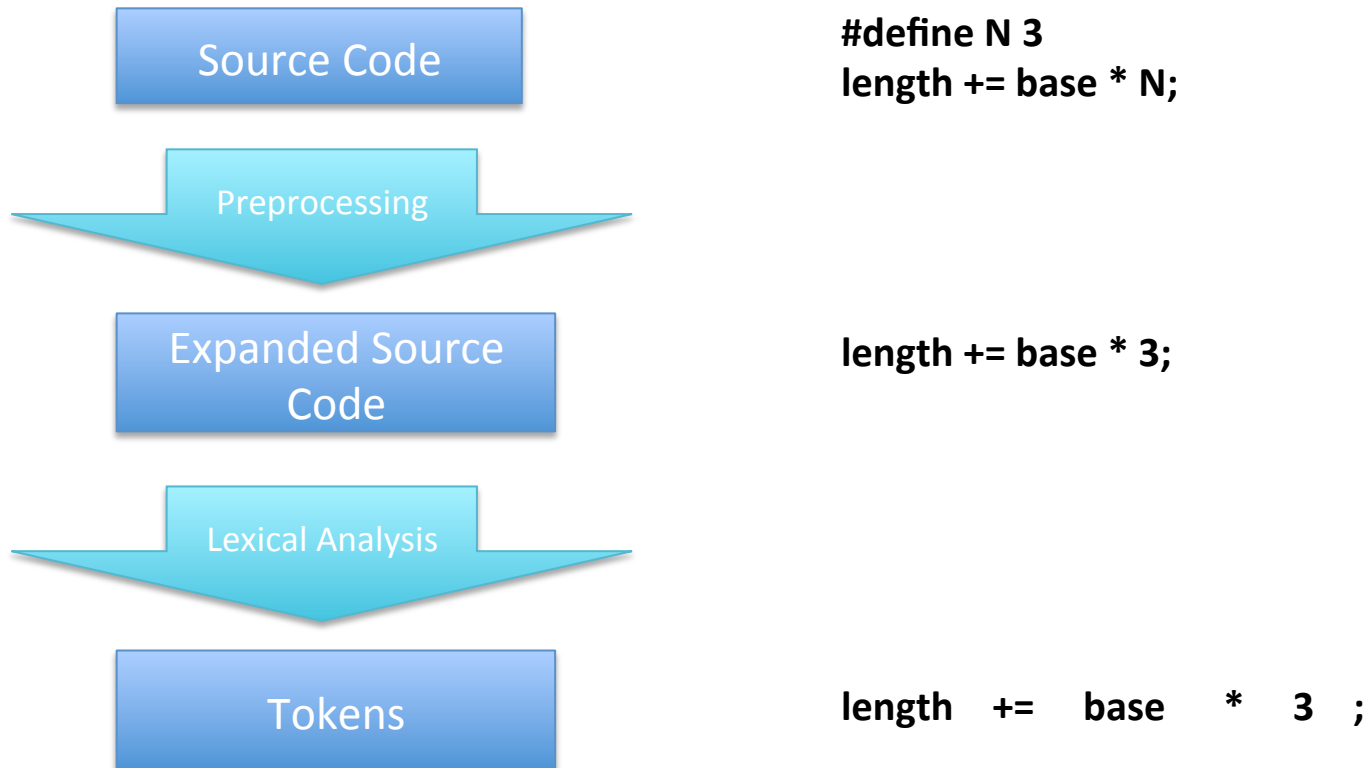
```
for ( int i = 0; i < SIZE; i++ )  
    ...
```



```
for ( int i = 0; i < 25; i++ )  
    ...
```

Steps in C Compilation

- Really, generating an executable has more steps than you might expect.




Steps in C Compilation

Tokens

`length += base * 3 ;`

Parsing

Parse Tree



Code Generation

Assembly Code


```
mov ebx, base
imul ebx, 3
iadd ebx, length
mov length, ebx
```

Assembling

Object Code

0010111010010101101

Steps in C Compilation



Looking for Tokens

- The compiler has to break the source into *tokens*
- This is called *lexical analysis*
- A token can be:
 - An *identifier* (e.g., a variable or a function name)
 - A *keyword* (e.g., void or while)
 - A *literal value* (e.g., 3.1415, or “Hello World”)
 - An *operator* (e.g., *, ++ or >=)
 - An *explicit separator* (e.g., (, } or ;)
- White space between the tokens is ignored
(except, of course, that it can separate tokens)

Identifiers

- *Identifiers* are variable, function or other names
- We get to choose names for these, but there are rules:
 - An identifier consists of **letters**, **digits** or **underscore**
 - But, a variable can't start with a digit
 - So **x2** is an identifier
 - but **2x** is a number followed by an identifier
- Identifiers are case sensitive in C
 - So **myValue** is not the same as **myvalue**

Reserved Keywords

- Of course, some words in C already have special meaning
- In C89, you can't use these as identifiers
auto, **break**, **case**, **char**, **const**,
continue, **default**, **do**, **double**, **else**,
enum, **extern**, **float**, **for**, **goto**, **if**,
int, **long**, **register**, **return**, **short**,
signed, **sizeof**, **static**, **struct**,
switch, **typedef**, **union**, **unsigned**,
void, **volatile**, **while**
- C99 adds a few more:
_Bool, **_Complex**, **_Imaginary**, **inline**,
restrict

Fun with Lexical Analysis

- What are the tokens in:
`a + ++b >= c-- -d`
- How about now:
`a+++b>=c---d`
- There are lots of things this could mean:

```
a ++ + b >= c - -- d    ?
```

```
a ++ +b > = c- - -d    ?
```

```
a + + + b > = c - - - d    ?
```

- This isn't really about precedence.
 - We can't even think about precedence until we know what the operators are.

The Scanner is Greedy

- The scanner works from left to right, grabbing the longest token it can
 - This is called *maximal munch*
- So, for our example:

a +++b>=c---d	(because a+ isn't a token)
a ++ +b>=c---d	(because +++ isn't a token)
a ++ + b>=c---d	(because +b isn't a token)
a ++ + b >=c---d	(because b> isn't a token)
a ++ + b >= c---d	(because >=c isn't a token)
a ++ + b >= c - ---d	(because c- isn't a token)
a ++ + b >= c -- -d	(because --- isn't a token)
a ++ + b >= c -- - d	(because -d isn't a token)

Be the Scanner

- In the following expression, how many tokens are there?

-----xy+ -=x++y* / z ;

- What are they?
- This expression wouldn't parse, but we can still talk about what the scanner would do with it.

Coding Style Conventions

- There are lots of ways you **can** write a working program
- But there's a difference between what you can do and what you should do
- Consider this submission from the first International Obfuscated C Code Contest:

```
int i;main(){for(;i["]<i;++i){--i;}"];read('-'-'-'-',i+++ "hell\
o, world!\n",'/'/'/'/')));}read(j,i,p){write(j/p+p,i---j,i/i);}
```

– I'm told it's a "Hello World" program

Coding Style Conventions

- Or this one from the 1993 contest:

```
05(02,07,03)char**07;{return!(02+==01+01)?00:!(02==02>01)?printf("\045\157\012"  
,05(012,07+01,00)):!(02==02>>01)?(**07<=067&&**07>057?05(03,07,*(*07)++-060+010  
*03):03)!:!(02--03--03)?(072>**  
07&&060<==**07?05(04,07,012*03-060  
+*(*07)++):03)!:!(02--!03+!!03)?(  
**07>057&&**07<=071?05(05,07,*(  
07)+++03*020-060):**07<=0106&&  
00101<==**07?05(05,07,020*03+*(*07)  
++-067):0140<**07&&**07<0147?05(05,  
07,-0127+*(*07)+++020*03):03)!:(  
02==02-01)?(*07==050?050**++*07,  
05(013,07,05(012,07,00)):*07<056  
&&054<*07?055**++*07,-05(06,07,  
00):054>**07&&052<**07?050*(*07)  
++,05(06,07,00)!:(*07^0170)||!(  
0130^**07)?*++*07,05(05,07,00):*  
*07==0144||**07==0104?++*07,05(04,  
07,00):05(03,07,00)):!--02?(*  
*07==052?05(07,07,03*)(**++*07,05(06  
,07,00))):!(045-**07)?05(07,07,  
03%(03+( *07)++,07,00))):!(**  
07^057)?05(07,07,03/*07,05(  
06,07,00))) :03)!:!(02++=01-02)?05(07  
,07,05(06,07,00)):!(02+==02/02)?(!(*  
*07-053)?05(011,07,03+(++*07,05(010,07,00))) :!(055^**07)?05(011,07,03-(03+*(*07  
)++,05(0010,07,00))) :03)!:!(02==0563&0215)?05(011,07,05(010,07,00)):(++*07,03);}
```


Coding Style Conventions

- These examples are fun (?) but deliberately hard to read and understand
- Normally, this is the opposite of what we want
- In CSC 230, we adopt some coding style conventions, rules for:
 - Naming conventions
 - Spacing and indentation
 - Where important comments go and what they contain
- Fortunately, editors can often help us with this.

CSC 230 Style Guidelines

- A Javadoc-style block comment at the top of each source file
 - With a `@file` tag giving the filename
 - And an `@author` tag with name and unity ID.
 - And a brief description of what the program does

```
/** I'm a program that reads in a list of words.  
    @file wordList.c  
    @author Bill Smith bsmith97  
    */
```

```
#include <stdio.h>  
#include <stdlib.h>
```

```
...
```

CSC 230 Style Guidelines

- A Javadoc-style block comment at the top of each function:
 - A sentence or two describing the function's purpose
 - @param tags for each parameter
 - @return tag for the return value

Omit if not needed.

```
/** Print a given number of copies of a given character.  
 * @param ch The character to print  
 * @param n The number of times to print ch.  
 */  
void printSymbol( char ch, int n )  
{  
    . . .
```

Extra stars are OK, if you like them.

CSC 230 Style Guidelines

- A good, Javadoc-style comment on each constant, global variable and type definition.
- Magic numbers, avoid bare constants for:
 - **Any value that could have an explanation**

```
area = radius * radius * 3.1415926;
```

- **Any potentially tunable parameter**

```
score += 350;
```

- **Any value that needs to occur multiple times**

```
for ( int i = 0; i <= 99; i++ )  
    ...;
```

```
for ( int j = 99; j >= 0; j++ )  
    ...;
```

CSC 230 Style Guidelines

- Curly bracket placement
 - For function definitions, it goes on the next line (to make functions stand out)
 - For everything else, it goes on the same line
- Indentation
 - No hard tabs, just indent with spaces (why?)
 - Indent using any number of spaces you want, 2 spaces, 3 spaces ... maybe 4 spaces.
 - But, be consistent.
- Just one statement per line

CSC 230 Style Guidelines

- Global Variables
 - Avoid these. Only use them where we tell you to.
 - So, functions communicate with the rest of the program via parameters and return values.

```
#include <stdio.h>
#include <stdlib.h>

/** Any code can access me. */
float cost = 25.88;

void someFunction()
{
    . . .
}
```

Don't do this,
unless the design
says to.