

# Arrays

CSC 230 : C and Software Tools  
NC State Department of Computer  
Science

# Topics for Today

- Declaring and using arrays
- Array initialization
- Operations on arrays
- Fixed- and variable-length arrays
- Arrays and sizeof()
- Multi-dimensional arrays
- Arrays and functions
- Strings as char arrays

# Making Arrays

- When you declare an array, you get to specify
  - The **type of its elements**
  - **How many elements** it will have
  - **How** they will be **initialized** (if at all)
- Example:  

```
int a[ 10 ];
```
- Notice, part of the type is given before the variable name, part **after**.
- In many contexts, the number of elements must be an integer **constant expression**

# Meet the Array

- Array notation is similar to Java
- You can access element  $i$  in array  $a$ :  $a[i]$ 
  - For an  $n$ -element array, indices range from  $0$  to  $n-1$
- You can also say it backward:  $i[a]$ 
  - Why can you do this?
  - It's a consequence of what the  $a[i]$  notation really means in C
  - We'll learn about this later

# C vs Java

- In C, declaring an array variable allocates the array.
- In Java, you need to declare and allocate an array.

```
int a[ 10 ];
```

- In C, arrays can be allocated statically, on the stack or on the heap (later).

```
int[] a = new int [ 10 ];
```

- In Java, all arrays are always allocated on the heap.

# Array Initialization

- Initialization rules are like other variables
  - If you don't initialize a **stack-allocated array**, it will contain whatever was previously in that memory.

```
int a[ 20 ];
for ( int i = 0; i < 20; i++ )
    printf( "%d ", a[ i ] );
```

outputs

0 1 1423550160 32767 182865920 ...

Yay, garbage!

- If you don't initialize a **statically-allocated array**, it starts out full of zeros.
- You can fill in the elements one at a time.

```
int a[ 20 ];
a[ 0 ] = 7;
a[ 1 ] = 14;
...;
```

# Array Initialization

- Or, you can ask the compiler to initialize it for you:

```
int a[ 5 ] = { 7, 14, 21, 28, 35 };
```

- If you initialize just some of the array, the rest will be initialized to zero.

```
int a[ 5 ] = { 7, 14, 21 };
```

Imagine ... , 0, 0

- This is called *partial initialization*.

# Array Initialization

- You can let the compiler determine the size
  - by measuring the length of your initialization list.

```
int a[] = { 7, 14, 21, 28, 35 };
```

OK. You must want  
a 5-element array.

- C99 offers an alternative initialization syntax

```
int a[] = { [7] = 10, [2] = 40, [5] = 35 };
```

0	0	40	0	0	35	0	10
---	---	----	---	---	----	---	----

# Array Initialization

- You can even mix the two notations
  - Where you give an index, C will put the given value there.
  - Where there's no index, C will fill the subsequent index.
  - If you don't give an array size, you'll get an array that's just large enough to hold the last value.

```
int a[] = { 1, [7] = 10, 90, [2] = 40, 2,  
           [5] = 35 };
```



# Array Bounds Checking

- C does **no bounds checking** for array access
  - Like many things in C, this decision favors performance over security
  - And, in many situations, the compiler may not even know the size of the array anyway.

```
int a[ 6 ] = { };
```

more program memory ... 0 0 0 0 0 0 ... more program memory

- What do you get if you index off the end?
  - Let's find out.

# Looking at Memory

- Example program `howdyNeighbor.c`

```
int x = 25;
int a[ 3 ] = { 10, 20, 30 };
int y = 50;

printf( "Array a: %d %d %d\n", a[ 0 ], a[ 1 ], a[ 2 ] );

printf( "A little before a: %d %d\n", a[ -2 ], a[ -1 ] );

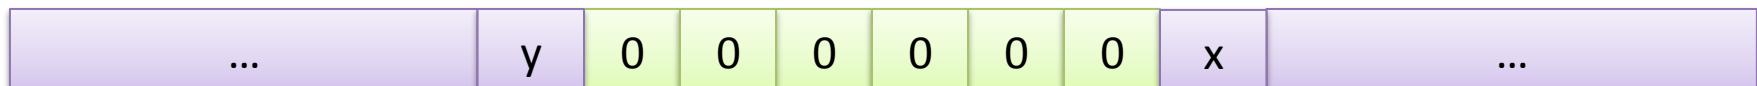
printf( "A little after a: %d %d\n", a[ 3 ], a[ 4 ] );
```

- Output:

```
Array a: 10 20 30
A little before a: 4195712 50
A little after a: 25 1876150896
```

# Looking at Memory

- From the example, I think this is what we're seeing:



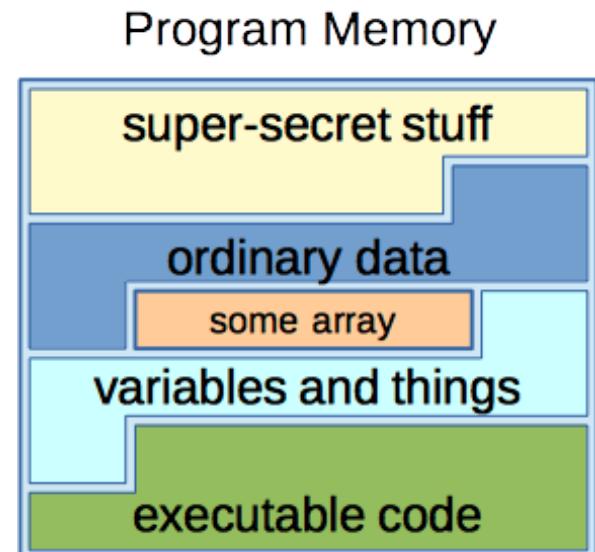
- In this example, we're looking around in other parts of the *stack frame*.
  - Where the function stores its local variables and other values it needs.
- So every array is like a little window into our program's memory

# Array Bounds Checking

- Accessing memory outside an array's bounds can do lots of neat things.
  - Accidentally overwrite some of your program's variables
  - Segmentation fault, bus error.
- These errors can be difficult to debug.
  - Maybe you overwrite an important variable or something your program needs to keep running.
  - If you're lucky, this may crash your program right away.
  - If you're unlucky, your program may crash ... later.

# With Great Power ...

- In Java, a bad index has defined behavior ... it will throw an exception.
- In C, behavior is **undefined** ... it could do anything ... probably nothing good ... for us.
- Accidentally writing past array bounds could be an exploitable vulnerability in our programs.
- An attacker could use this to change the data or behavior of a program.



# Buffer Overflow

- A *buffer overflow* happens when a program writes into memory past the end of an array.
  - Typically while reading input from a user or another program.
- It's an easy mistake to make:

```
printf( "Please enter some numbers\n" );  
  
int list[ 20 ];  
int i;  
while ( scanf( "%d", &list[ i ] ) == 1 )  
    i++;
```

# Operations on Arrays

- Normally, we work with an array one element at a time.
- C syntax includes very little for working with the whole array at once.
  - You can ask the size of an array:

```
printf( "%zd\n", sizeof( list ) );
```

- But you **can't** copy from one array to another:

```
list1 = list2; X
```

- Or concatenate two arrays, etc.

# Operations on Arrays

- Instead we typically need **loops** to process everything in an array:

```
int arrayCopy[ SIZE ];  
  
for ( int j = 0; j < SIZE; j++ ) {  
    arrayCopy[ j ] = arraySource[ j ];  
}
```

# Choosing Array Size

- The size of a statically allocated array must be a constant expression.

```
int sequence[ 100 ];  
  
int main()  
{  
    ...  
}
```

- This is true in some other contexts also.
- ... and used to be true for stack-allocated arrays.

# Variable-Length Arrays

- C99 permits the size of stack-allocated arrays to be determined during execution.

```
int main( void )
{
    int n;

    printf( "How many values would you like: " );
    scanf( "%d", &n );

    int list[ n ];
    ...
}
```

This is called a *variable-length array*.

What if this is zero?

or negative?

# Variable-Length Arrays

- Variable-length arrays aren't resizable
  - Once you choose the size, you're stuck with it for the array's lifetime.

```
int n;  
printf( "How many values? " );  
scanf( "%d", &n );  
int list[ n ];
```

*... do some stuff...*

```
printf( "Now how many? " );  
scanf( "%d", &n );  
int list[ n ];
```

This won't work; looks like  
you're trying to declare a  
new array.

# Variable-Length Arrays

- But, you can choose a different size every time the variable goes into scope:

```
for ( int i = 0; i < 10; i++ ) {  
    int n;  
    printf( "How many values? " );  
    scanf( "%d", &n );  
    int list[ n ];  
    ...;  
}
```

This will work.

# Determining Array Size

- You remember the `sizeof` operator ... right?
- It reports the size of its argument, including array arguments.
  - An array's size is the number of elements times the size of each element.

```
#include <stdlib.h>

int a;
float b[100];
double c[ 16 ];

size_t aSize = sizeof( a );
size_t bSize = sizeof( b );
size_t cSize = sizeof( c );
```

Three variables

How much memory does  
each one take?

# Computing Array Size

- The `sizeof` operator can let you figure out array length (number of elements).

```
float b[ 100 ];  
...  
int len;  
len = sizeof( b ) / sizeof( b[0] );
```

It may not look like it, but this is really a constant expression.

Why not just say 100 instead?

This is a great trick ... but it doesn't work everywhere.

- Here, `sizeof()` can be evaluated at compile time.

# Multi-Dimensional Arrays

- To make a multi-dimensional array, you stack up square brackets after the name.
- Declared like this, you have to give a size for all dimensions.

```
int table[ 3 ][ 4 ];
```

Outer dimension  
(normally we'd think of this as rows)

Inner dimension  
(normally we'd think of this as columns)

- Accessing an element, just like you'd expect,
  - an index for row, then column, inside square brackets.

```
x = table[ 2 ][ 3 ];  
table[ 1 ][ 0 ] = 99;
```

# Multi-Dimensional Array Initialization

- With initialization, elements are initialized
  - row-by row, each row in curly brackets
  - and, left-to-right within each row.

0	1	4	9
16	25	36	49
64	81	100	121

```
int table[ 3 ][ 4 ] = {  
    {0, 1, 4, 9},  
    {16, 25, 36, 49},  
    {64, 81, 100, 121 }  
};  
  
for( int i = 0; i < 3; i++ )  
    for( int j = 0; j < 4; j++ )  
        printf( "%d ", table[i][j] );
```

I'm the first  
row.

I'm the first  
element on  
the first row.

Prints 0 1 4 9 16 ...

# Partial Initialization

- You can do partial initialization in a 2D array

```
int table[3][4] = {  
    { 0, 1 },  
    { 16, 25, 36, 49 },  
    { 64 }  
};
```

You're going to  
get some zeros.

0	1	0	0
16	25	36	49
64	0	0	0

- Or, with C99 syntax

```
int table[3][4] = {  
    [ 1 ] = { 16, 25, 36, 49 },  
    [ 0 ] = { [ 1 ] = 1 },  
    [ 2 ][ 0 ] = 64  
};
```

0	1	0	0
16	25	36	49
64	0	0	0

# Multi-Dimensional Array Initialization

- You can let the compiler determine the array size
  - But just the outermost (first) dimension.

```
int a[ ] [ 3 ] = { { 10, 15, 20 },  
                    { 30, 35, 40 } };
```

OK

```
int a[ 2 ] [ ] = { { 10, 15, 20 },  
                    { 30, 35, 40 } };
```

Not OK

- Why the difference?
  - Size of all inner dimensions must be specified in the declaration.
  - It's used by the compiler to write the indexing code.

# Memory Organization

- In memory, two-dimensional arrays are laid out in *row-major order*

0	1	4	9
16	25	36	49
64	81	100	121

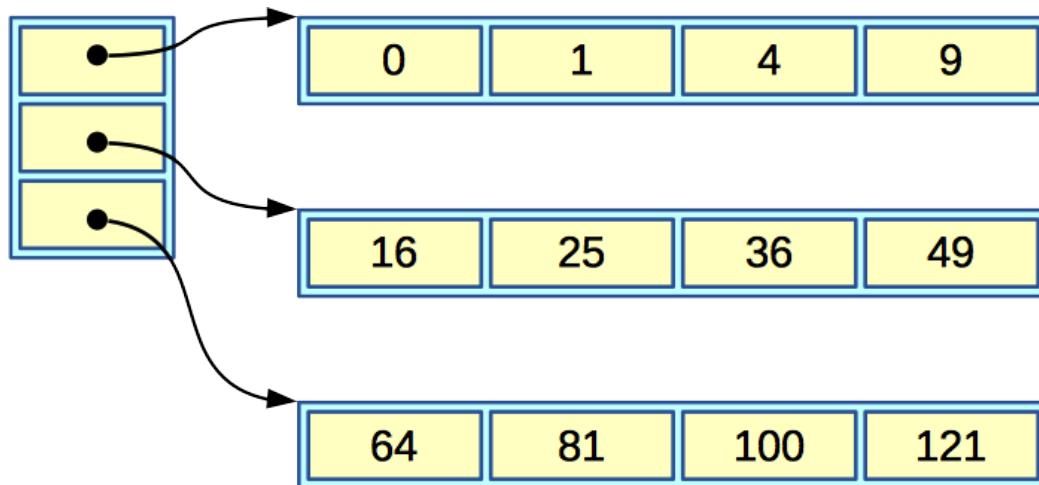
You can think  
of it like this.

0	1	4	9	16	25	36	49	64	81	100	121
---	---	---	---	----	----	----	----	----	----	-----	-----

But, in memory  
it's like this.

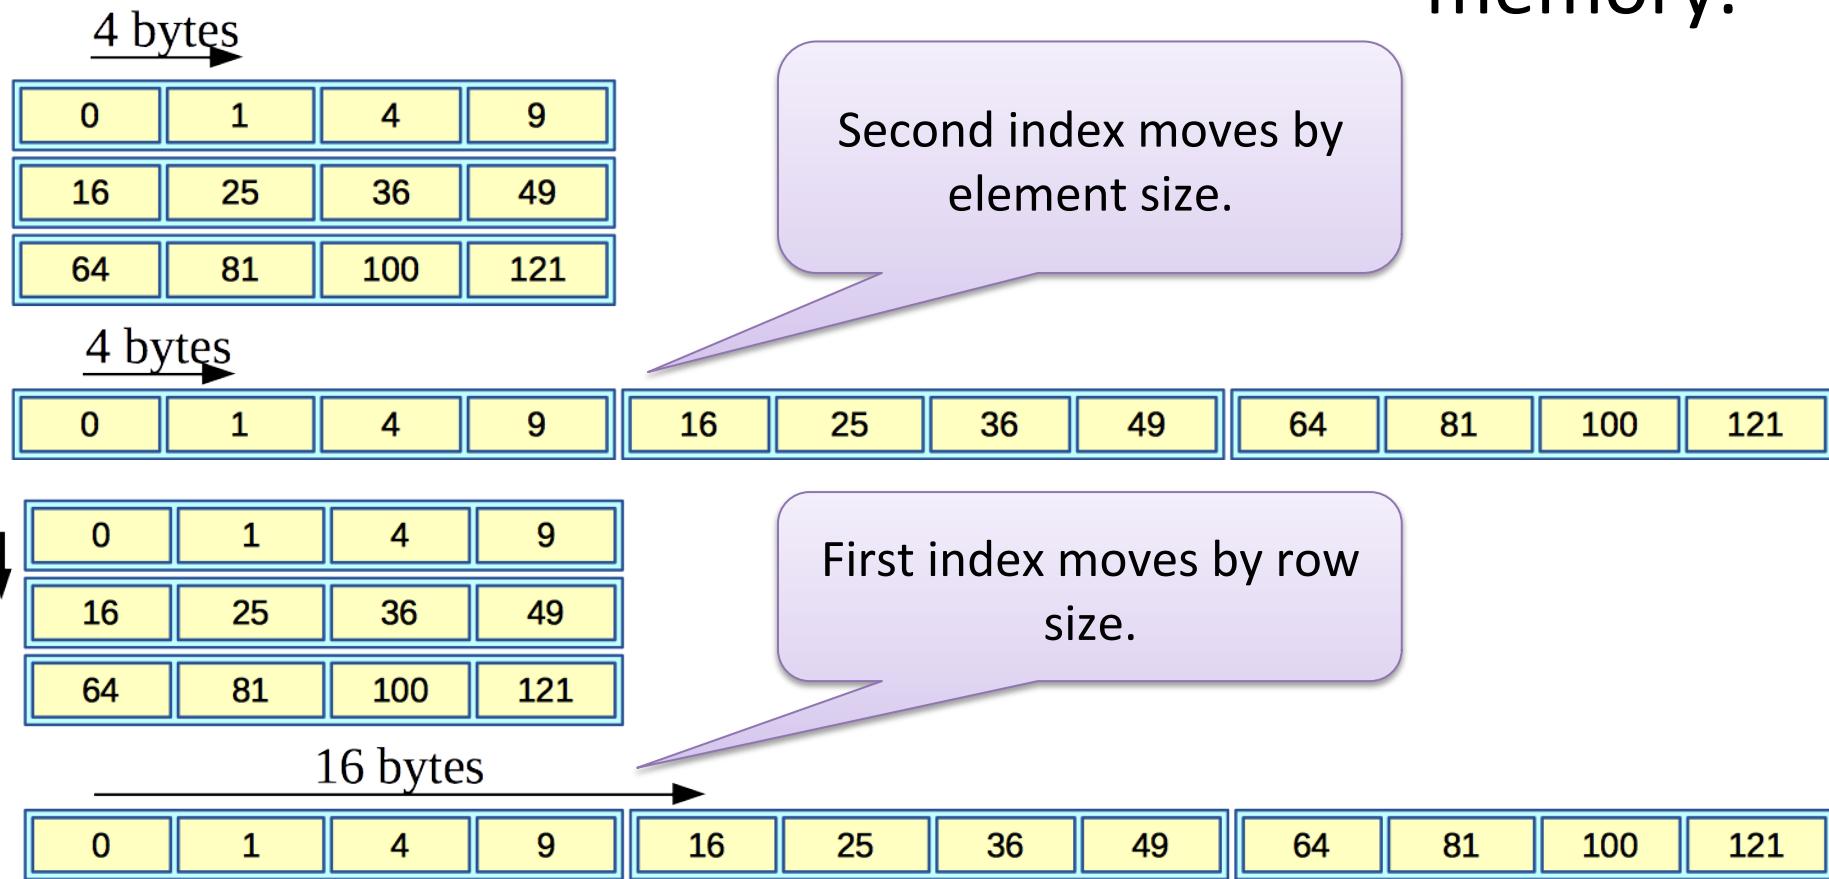
# Memory Organization

- It's not like this in Java.
  - It was more like this:



# Memory Organization

- Fixed row size makes it easy to index into this memory.



# Multi-Dimensional Arrays

- Multi-dimensional arrays are just really **arrays of arrays**.

```
int row[ 4 ];
```

I'm an array of 4 ints.

```
int table[ 3 ][ 4 ];
```

I'm an array of 3 guys  
like row.

- We can take an array apart, into its individual elements (rows)

- .. and each row's elements (ints) I'm the size of the whole array.

```
sizeof( table );
```

I'm the size of one of its  
elements, a row.

```
sizeof( table[ 0 ] );
```

```
sizeof( table[ 0 ][ 0 ] );
```

I'm the size of one of that row's  
elements.

# Arrays and Stack Space

- Arrays can be stack-allocated.
- But, in practice, stack space is limited.
- A big array on the stack is an easy way to use it up.
- Or, deeply nested recursion will do it too.

I'm probably going to segfault.

```
int bill()
{
    int smallTable[ 10 ];

    bigTable[ 0 ] = 1;
    return bigTable[ 0 ];
}
```

```
int bill()
{
    int bigTable[ 100000000 ];

    bigTable[ 0 ] = 1;
    return bigTable[ 0 ];
}
```

# Allocating Large Arrays

- Larger arrays can be allocated statically
  - The compiler can plan ahead for static storage.

The compiler will be sure  
to leave room for this.

```
int bigTable[ 100000000 ];  
  
int bill()  
{  
    bigTable[ 0 ] = 1;  
    return bigTable[ 0 ];  
}
```

- Or a large array could be dynamically allocated  
... more about this later.

# Arrays as Function Parameters

- An array can be passed as a function parameter
- The array parameter looks like an array declaration
- You can specify the array length **explicitly** in the function declaration ... if you want.

```
void report( int table[8] )  
{  
    ...  
}
```

Array parameter,  
with explicit size.

```
void report( int table[5][8] )  
{  
    ...  
}
```

# Array Parameter Dimensions

- In array passing, the compiler doesn't care so much about the outermost size.
- But it cares a **lot** about the inner sizes.

```
void report( int table[5][8] )  
{  
    ...  
}
```

This will work

```
int regularTable[5][8];  
int tallTable[6][8];  
int wideTable[5][9];  
...  
report( regularTable );  
report( tallTable );  
report( wideTable );
```

So will this.

But not this!

# Array Parameter Dimensions

- In array indexing, the inner size determines the width of each row.
  - The compiler needs to know this to index between the rows..
- But, the outermost index just determines how far we go before we're out of bounds
  - And C doesn't do any bounds checking anyway.

The compiler will let you omit this.

And most people do, since it doesn't matter anyway.

```
void report(int table[][12])
```

Same here.

```
{  
    float analyze(float list[])
```

Same here.

```
    {  
        ...  
    }
```

```
        int process(short block[][7][10])
```

```
        {  
            ...  
        }
```

# Array Parameter Dimensions

- The function will work for any array length (at least for the outer index).
- It can't even tell how big the array is.

```
void report(int list[])
{
    int items = sizeof(list)/sizeof(list[0]));
}
```



Size of the whole array ...  
divided by size of an element ...  
should be number of elements.

But, the function  
doesn't know the  
size of the whole  
array.

This compiles, but it  
doesn't mean what  
you want it to.

# Communicating Array Size

- How can a function tell how big the array is?
- Often, we need to pass that in as another parameter.

```
int countValues(int len, int list[])
{ ... }
```

- For multi-dimensional, fixed-width arrays, this usually just applies to the outer-most dimension.

```
int countValues(int rows, int table[][][7])
{ ... }
```

- In some special cases, the function may be able to figure out how long the array is.

```
void countChars(char str[])
{ ... }
```

# Passing Variable-Sized Arrays

- For variable-sized arrays, you **must** pass everything but the outer-most dimension.

You need this just to pass the array.

See, it gets used here in the type.

```
void sumTable(int rows, int cols, int table[][]cols)
{
    ...
    for ( i = 0 ; i < rows; i++)
        for ( j = 0; j < cols; j++)
            total += table[i][j];
    ...
}
```

You may also need this to know how big the array is.

# Passing Variable-Sized Arrays

```
void sumTable(int rows, int cols, int table[][][cols])
{
    ...
}
```

- We depend on the caller to tell us the truth.

```
int tbl[10][12];

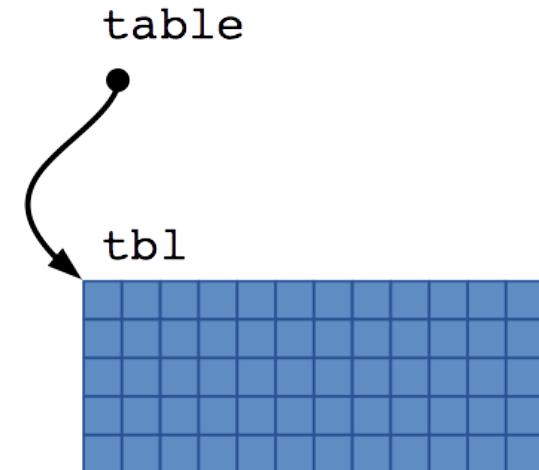
// Good client code.
sumTable( 10, 12, tbl );
// Evil client code.
sumTable( 20, 13, tbl );
```

# Arrays as Parameters

- Secretly, arrays are automatically **passed by reference (by starting address)**, not by value
  - This is why `sizeof()` doesn't do what you think it should for an array parameter

```
void useTable( int table[5][12] )
{
    size_t sz = sizeof( table );
}
```

```
int tbl[5][12] = { { ... } };
...
useTable( tbl );
```

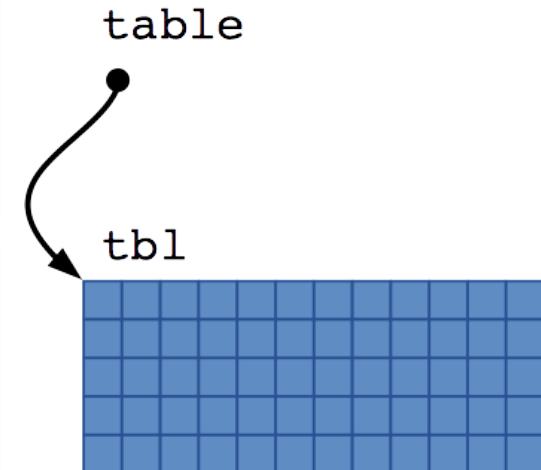


# Arrays as Parameters

- So, the function **can modify** the contents of the caller's array

```
void changeTable( int table[5][12] )
{
    table[3][6] += 1;
}
```

```
int tbl[5][12] = { { ... } };
...
changeTable( tbl );
```



# Array Return Values

- Functions can have array parameters
- ... but not array return values.

```
int makeList( int a, int b )[ 10 ]  
{  
    ...  
}
```

I guess it would look like this, if you could do it ... which you can't.

- Later, we'll figure out how to approximate this
  - by returning a pointer to an array.

# Character Strings

- In C, a **string** is:
  - A sequence of characters stored in a one-dimensional array.
  - With a null character to mark the end.
- You can still create plain-old character arrays:

```
char a[] = { 'c', 's', 'c', '2', '3', '0' };
```

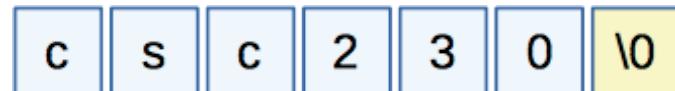
sizeof( a ) will report 6.



- You also have a more convenient syntax for creating null-terminated strings:

```
char s[] = "csc230";
```

sizeof( s ) will report 7.



# Literal String Concatenation

- String concatenation works in the initialization syntax.

```
char s1[] = "Now " "is " "the " "time";  
printf( "%zd\n", sizeof( s1 ) );
```

I'm 16 bytes.

- But, you can only use this on literal strings.
  - It doesn't do general-purpose string concatenation.

```
char s3[] = s s1;
```

This won't work.

# Strings in C

- Null termination is a convention in C
  - It lets us mark the end of a string in memory, without explicitly storing the length.
  - Standard string processing functions expect this.
  - For example, `printf()` can print a string with `%s`
  - ... provided it's null terminated.

```
char word[] = "csc203";
printf( "The string is %s\n", word );
```

# Reading Strings

- `scanf()` will read strings with `%s`
  - It reads a single word (a space-delimited sequence of non-whitespace characters)
- `Scanf` follows the convention of null-terminating the string it reads.

I have room for a 29-character string.

```
char word[ 30 ];  
scanf( "%s", word ),
```

Notice, no &.

Look, a buffer-overflow vulnerability.