

Dynamic Memory Allocation

CSC 230 : C and Software Tools
NC State Department of Computer
Science

Topics For Today

- Void pointers
- The free memory heap
- Using malloc/free
- Dynamic allocation hazards
- Copying strings
- Resizable arrays
- Memory performance comparison

Void Pointers

- You can't have a value of type void

```
X void v;
```

- But you can make a pointer to void

```
void *vPtr;
```

- And that void pointer can hold any pointer value you want

```
float f = 3.14;  
vPtr = &f;
```

You don't even
need a cast.

```
long l = 123456789;  
vPtr = &l;
```

Void Pointers

- But, you can't directly dereference a void pointer.

```
printf( "%ld\n", *vPtr );
```

- Think of a void pointer as a generic way of holding an address
 - they could be used to point to anything
 - but, to use what they point to, you need to know its type.

Using Void Pointers

- You need to convert a void * to a more specific type before you dereference.

```
long *lPtr = vPtr;  
printf( "%ld\n", *lPtr );
```

You don't even need
a cast.

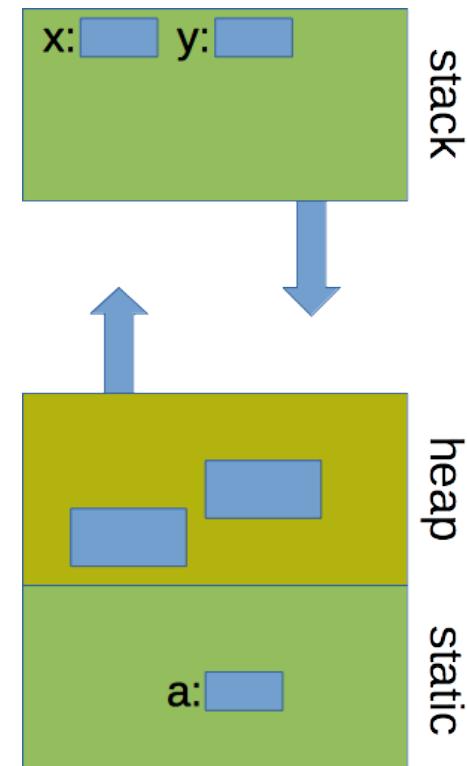
```
lPtr = (long *)vPtr;  
printf( "%ld\n", *lPtr );
```

But, it's probably a
good idea.

- Void pointers are used:
 - As parameter/return types for functions that can work with memory (no matter what it stores)
 - As temporary holders for pointers of various types.

Another Type of Memory

- We already have two types of memory
 - Stack space : for values with a lifetime tied to a single function execution
 - Static space : for values with a lifetime from the start to the end of the program's execution
- But, there's another type of memory, the **heap**
 - The program says when it needs heap memory (and how much)
 - The program says when it's done with a particular value in heap memory
 - Called: **dynamic memory allocation**



Asking for Heap Memory

- First, you must include `stdlib.h`

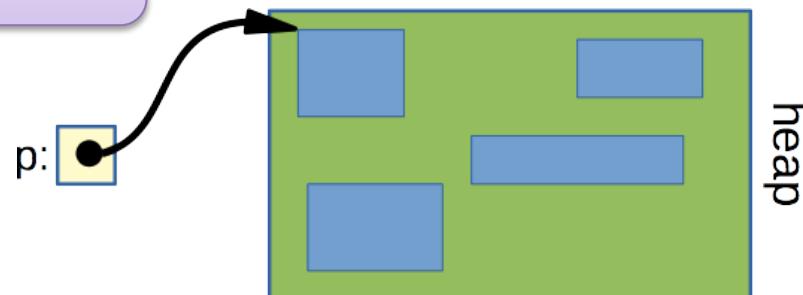
```
#include <stdlib.h>
```

- The C Standard Library has functions to let us request heap memory

```
void *malloc( size_t size );
```

Here's a pointer
to your memory.

How many bytes
do you want?



Using malloc()

- To use malloc():
 - Pass in the number of bytes you need, probably using sizeof()
 - Store the result in a pointer for the type you need.
- Allocate a 1000-element array of ints:

```
int *list = (int *) malloc( 1000 * sizeof(int) );
```

This cast is a good idea,
but a C compiler doesn't care.

- Allocate space for a copy of a string:

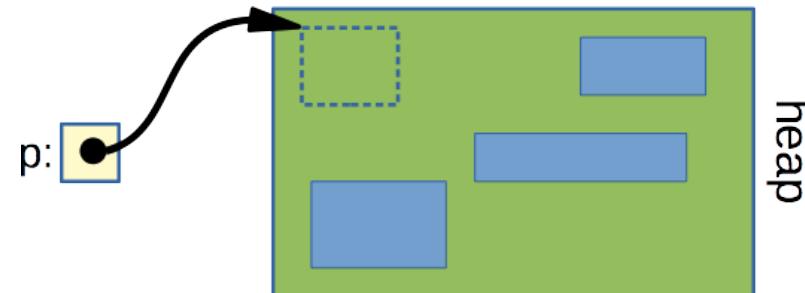
```
char *str2 = (char *) malloc( strlen( str ) + 1 );
```

Freeing Memory

- You get to keep each block of dynamically allocated memory as long as you need.
- You must tell the C standard library when you are done with a block:

```
void free( void *ptr );
```

Pointer to previously allocated
memory.



Using Malloc/Free

- Typical usage:

Allocate storage for the thing you need.

```
int *list = (int *)malloc(1000 * sizeof(int));
```

```
... list[ 5 ] = b; ...
```

Use this thing, probably across many ...

```
... z = list[ i ]; ...
```

... different ...

```
... list[x] = list[y];
```

... functions.

```
free(list);
```

Free the storage, probably in some other part of your program.

A Trivial Example

- A 20-element array

Get a 20-element array of doubles.

```
double *list = (double *)malloc(20 * sizeof(double));  
  
for ( int i = 0; i < 20; i++ )  
    list[ i ] = i * 1.5;  
  
for ( int i = 0; i < 20; i++ )  
    printf( "%d : %.2f\n", i, list[ i ] );  
  
free( list );
```

Use it.

Free it.

A Better Example

- Read and return a list of integers.

Here's how we return the array.

Here's how we return the length.

```
int *readList( int *len )
{
    scanf( "%d", len );

    int *list = (int *)malloc( *len * sizeof( int ) );

    for ( int i = 0; i < *len; i++ )
        scanf( "%d", list + i );

    return list;
}
```

A Better Example

- Code to use `readList()`

```
int len;  
int *list = readList( &len );  
  
for ( int i = 0; i < len; i++ )  
    printf( "%d\n", list[ i ] );  
  
free( list );
```

Get the list.

Use it.

Free it when we're
done.

Using Dynamically Allocated Memory

```
void *malloc( size_t size );
```

```
void free( void *ptr );
```

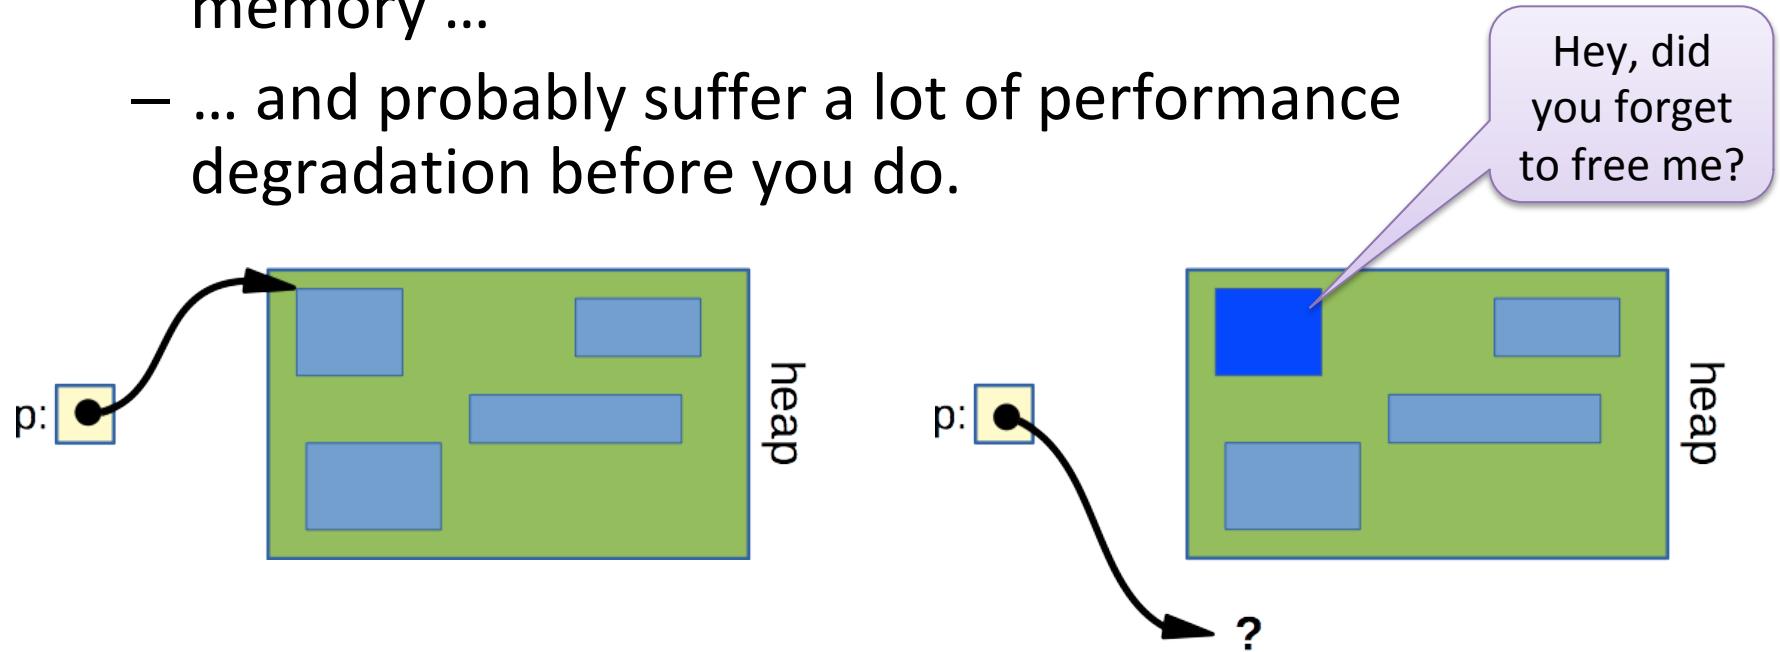
- Memory returned by malloc() is **uninitialized**, it may contain garbage from a previous allocation
- Returns **NULL** on failure (e.g., if memory is full)
- Type `size_t` for representing the size of storage regions
 - It's what `sizeof()` evaluates to
 - And what `printf()` prints with format string "%zd"
 - Just another name for an integer type (unsigned long on the common platform)

Dynamic Allocation Hazards

- All previous warnings apply
 - Buffer overflow, we can accidentally read/write off the end of a dynamically allocated array.
 - Dereferencing a NULL or uninitialized pointer
- But, now there are some new ways to make mistakes.

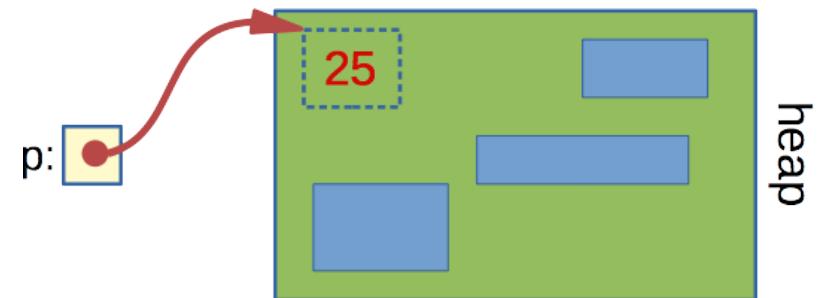
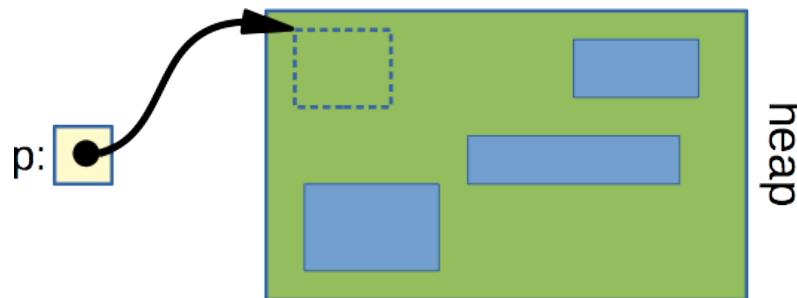
Leaking Memory

- What if you forget to free memory when you're done with it?
 - That's called a *memory leak*.
 - If you keep doing this, you'll eventually run out of memory ...
 - ... and probably suffer a lot of performance degradation before you do.



Dangling Pointers

- Freeing memory doesn't make your pointer variable go away (or set it to NULL)
- So, there's a possibility of using a pointer to memory you've already freed.
 - This is called an *dangling pointer*
 - If you're lucky, using a dangling pointer will crash your program right away.
 - If you're unlucky, ... well maybe much later.



Using Dynamically Allocated Memory

- Other, less common mistakes:
 - Free memory you didn't malloc.

```
int *p = (int []) { 5, 10, 15, 20 };  
...  
free( p );
```

- Free the same memory twice.

```
char *p = (char *)malloc( 1024 );  
...  
free( p );  
...  
free( p );
```

Thinking about Dynamic Allocation

- In Java, we have dynamic memory allocation; that's what `new` does.
 - In fact, we usually have to dynamically allocate anything larger than a primitive type (`int`, `double`, `char`, etc.)
 - But, Java performs **Garbage Collection**; it automatically finds memory we can no longer use and reclaims it as free space.
- Standard C doesn't perform garbage collection
- We need to tell the system when we're done with each block

Thinking about Dynamic Allocation

- Why use malloc/free?
 - When you need to store something that must live longer than a single function call.
 - When you need something with size that's not known until execution.
 - When you need to store something that's too large to put on the stack.
- But, only use it where you need it.
 - Stack and static storage can be much less expensive to allocate/free.
 - And, they may have less overhead to use.
 - And, they are easier to work with.

Saving Strings

- Consider this (bad) parsing code, to read a list of strings.

```
int readWords( char *words[],
               int capacity )
{
    int count = 0;

    char str[ 100 ];
    while ( count < capacity &&
            scanf("%99s", str) == 1 ) {
        words[ count ] = str;
        count++;
    }

    return count;
}
```

Where to store the list of strings.

Keep reading strings as long as you can.

Save them in the array.

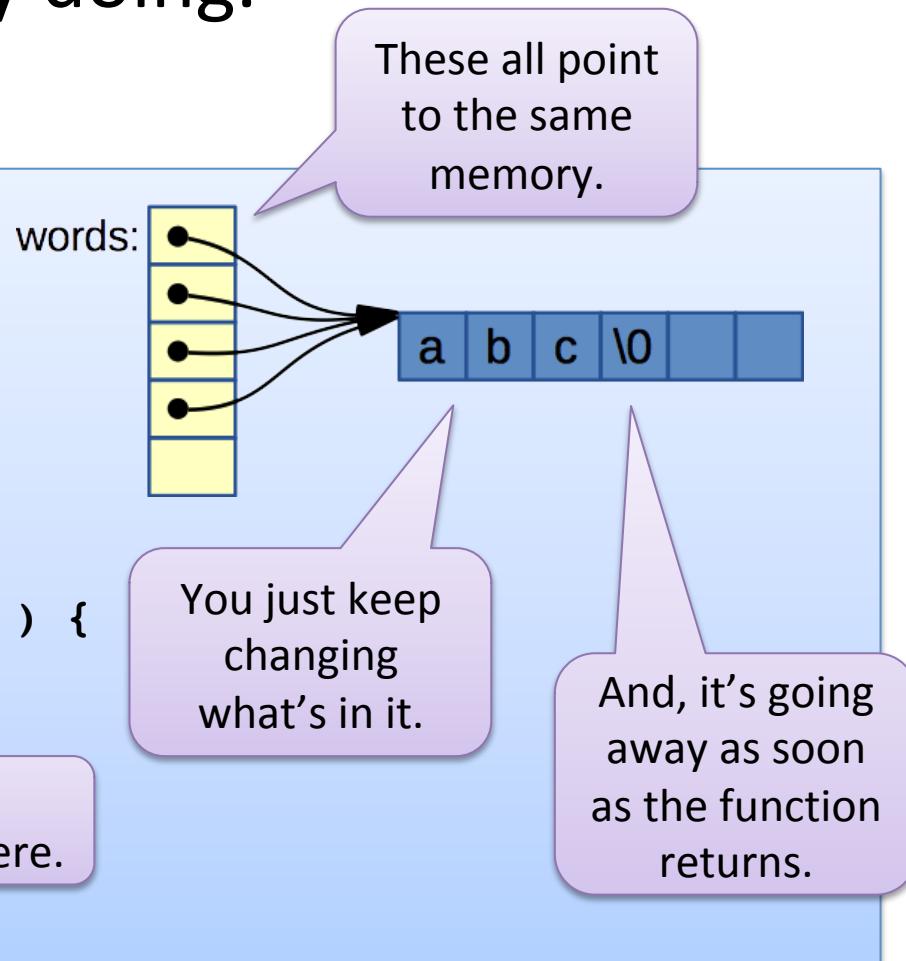
Saving Strings

- What this code is really doing.

```
int readWords( char *words[],
                int capacity )
{
    int count = 0;

    char str[ 100 ];
    while ( count < capacity &&
            scanf("%99s", str) == 1 ) {
        words[ count ] = str;
        count++;
    }

    return count;
}
```

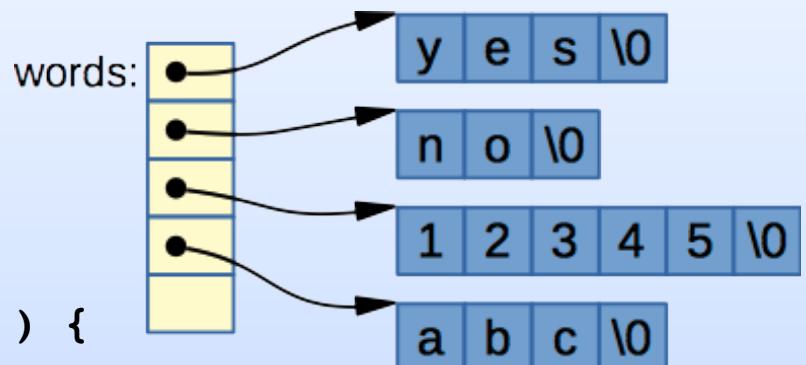


Saving Strings

- This code is better. It copies each string to its own memory.

```
int readWords( char *words[],
               int capacity )
{
    int count = 0;
    char str[ 100 ];
    while ( count < capacity &&
            scanf( "%99s", str ) == 1 ) {
        words[ count ] =
            (char *) malloc( strlen(str) + 1 );
        strcpy( words[ count ], str );
        count++;
    }

    return count;
}
```



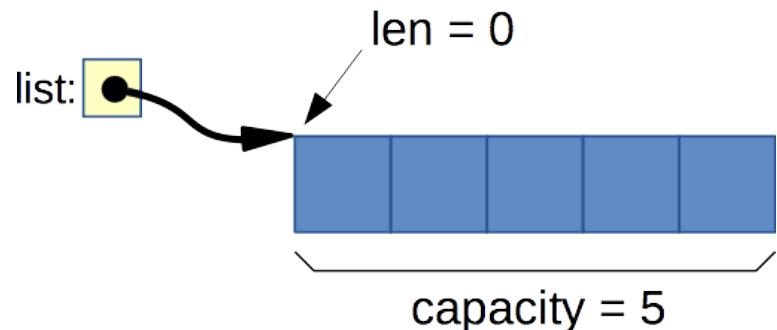
This memory
will stay
around until
you free it.

Resizing Arrays

- Our earlier example, `readList()`, was easy
 - We were told how many items to expect, so we could make our array exactly the right size.
- Imagine we had to read a list of values, without knowing *a priori* how many there were
 - We could use a dynamically allocated array ...
 - ... but we'd need to be able to enlarge it if we ran out of room.
 - This is what a **resizable array** does

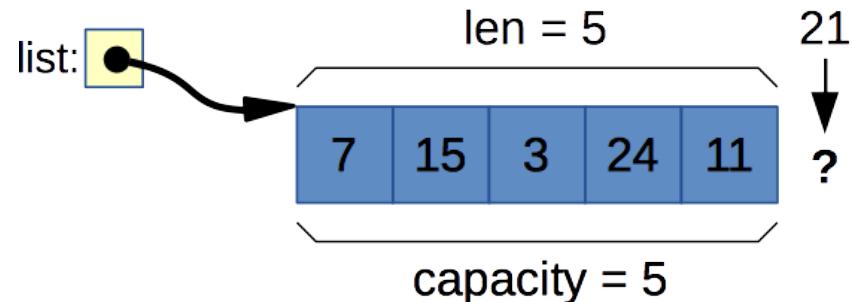
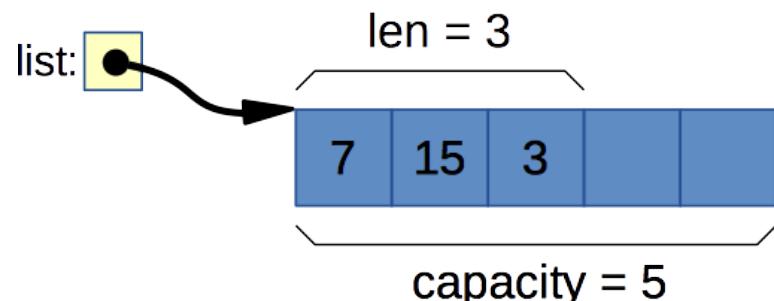
Resizing Arrays

- We need to keep up with three things
 - A pointer to our array (list)
 - The number of elements used (len)
 - The capacity of the array (capacity)
- Start out with an array that's large enough ... for now.



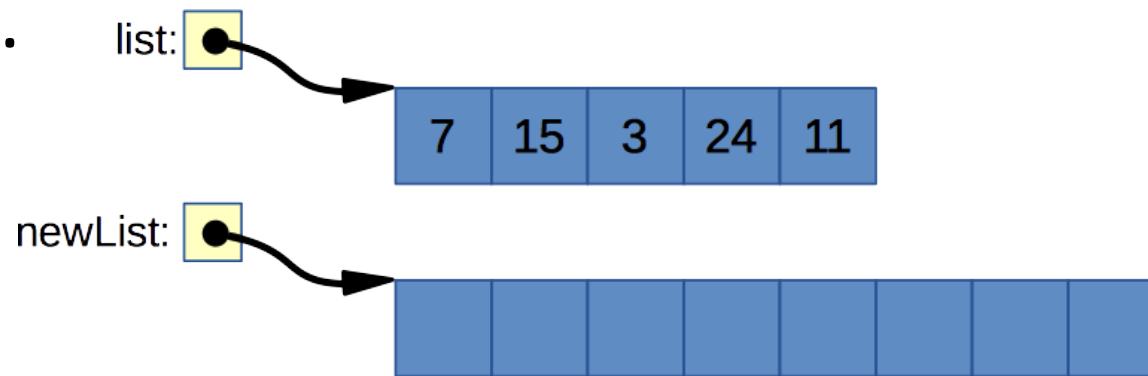
Resizing Arrays

- As we append items:
 - len tells us where to put the next one.
 - increment len on each append
- Everything's easy until until we reach the capacity.

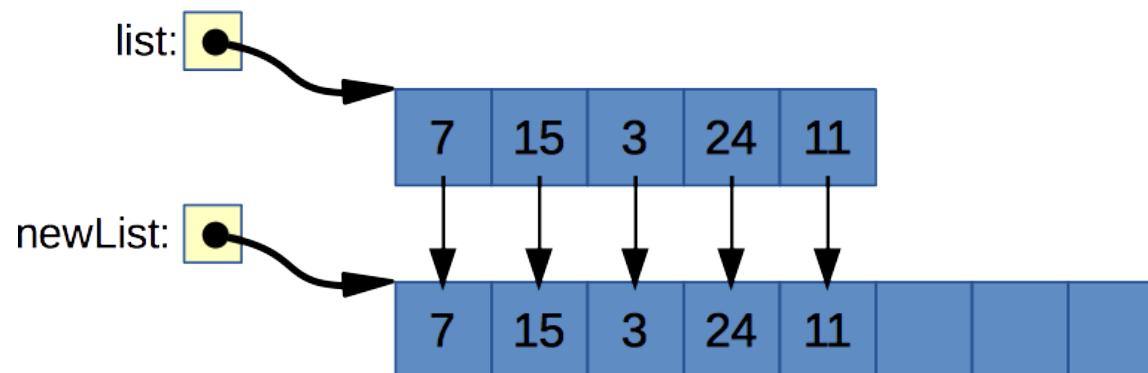


Resizing Arrays

- That's OK. We'll just allocate a new, larger array.

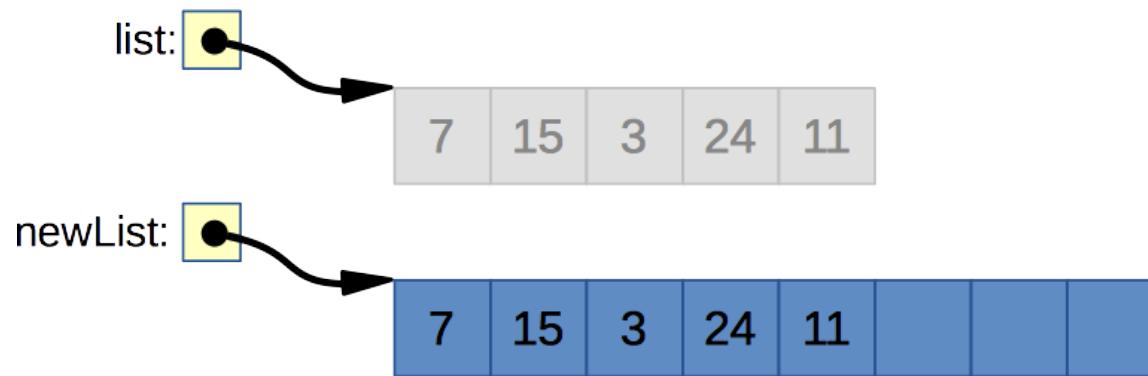


- Then copy everything over to the new array.

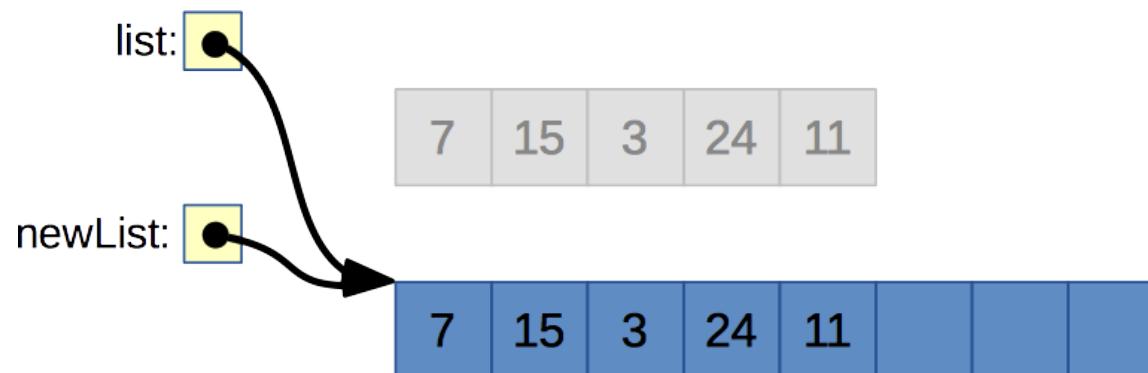


Resizing Arrays

- Then, we can free the old array.

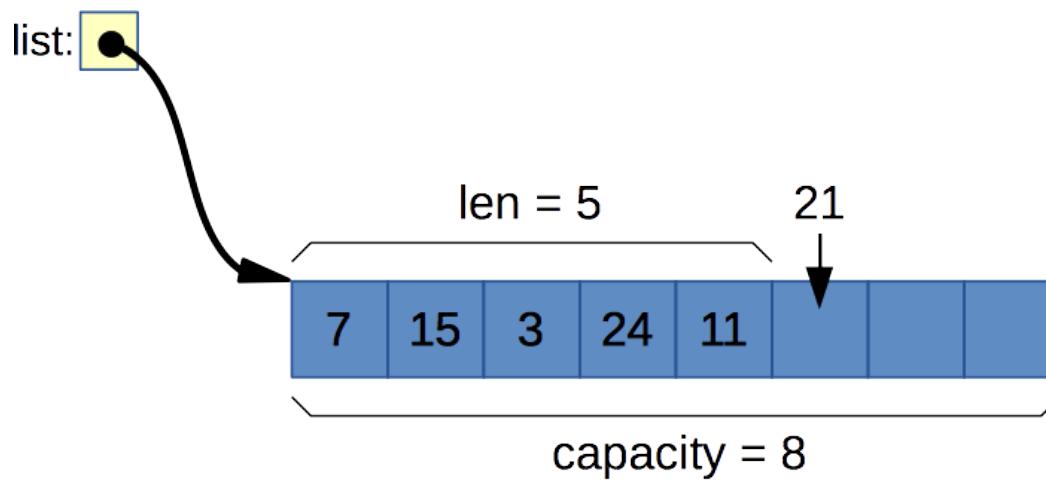


- And start using the new array instead.



Resizing Arrays

- Now, you have more room to insert new items.



- There are a few example programs that show this technique.

Resizing an Array, the Slow Way

```
int capacity = 5;
int len = 0;
int *list = (int *)malloc( capacity * sizeof( int ) );

int val;
while ( scanf( "%d", &val ) == 1 ) {
    if ( len >= capacity ) {
        capacity += 1;
        int *newList = (int *)malloc(capacity * sizeof(int));
        for ( int i = 0; i < len; i++ )
            newList[ i ] = list[ i ];
        free( list );
        list = newList;
    }
    list[ len++ ] = val;
}
```

Bad Idea, just grow a little

Allocate larger array

Copy everything over

Switch to the new array

Copying Memory

- The C Standard Library has functions to copy blocks of memory.

```
void *memcpy( void *dest, void const *src, size_t n );
```

from here.

to here.

Copy this many bytes.

- Here, src and dest blocks can't overlap.
- But they can with memcpy's slightly more expensive friend:

```
void *memmove( void *dest, void const *src, size_t n );
```

A Better Way

```
int capacity = 5;
int len = 0;
int *list = (int *)malloc( capacity * sizeof( int ) );

int val;
while ( scanf( "%d", &val ) == 1 ) {
    if ( len >= capacity ) {
        capacity *= 2;
        int *newList = (int *)malloc(capacity * sizeof(int));

        memcpy( newList, list, len * sizeof( int ) );

        free( list );
        list = newList;
    }

    list[ len++ ] = val;
}
```

That's better.

Help from a standard function to copy everything over.

More storage overhead, but much less time overhead.

Meet realloc()

- There's a C standard library function to help.

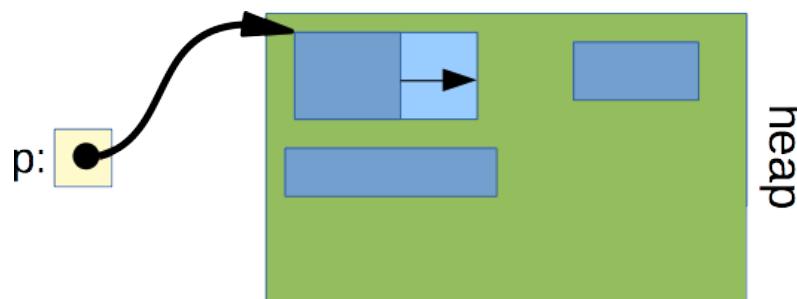
```
void *realloc(void *ptr, size_t size);
```

Pointer to larger
block.

Pointer to previously
allocated memory.

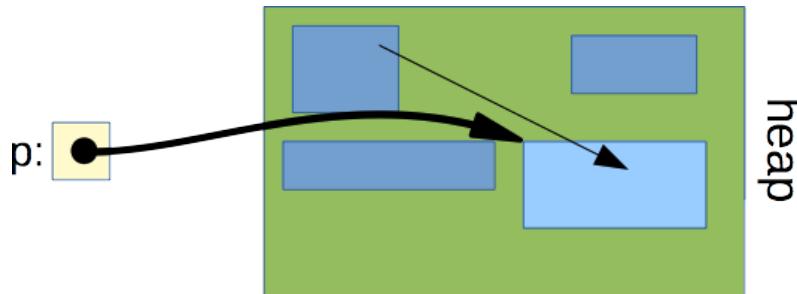
Size you need now.

- realloc() will enlarge your block and give you back the same pointer ... if it can.



Meet realloc()

- If it can't enlarge your block, realloc() will:
 - find a larger block
 - copy everything over
 - free the old block
 - return a pointer to the new block.
- This will let us simplify our resizable array code.



Using Realloc

```
int capacity = 5;
int len = 0;
int *list = (int *)malloc( capacity * sizeof( int ) );

int val;
while ( scanf( "%d", &val ) == 1 ) {
    if ( len >= capacity ) {
        capacity *= 2;
        list = (int *)realloc( list, capacity * sizeof( int ) );
    }

    list[ len++ ] = val;
}
```

Choose new capacity.

Enlarge array and start using it.

Beware, this isn't going to work so well if we actually run out of memory.

Meet calloc()

- There's one more dynamic allocation function

```
void *calloc(size_t count, size_t size);
```

Pointer to your memory.

Number of items you need.

Size of each item.

- Good for allocating arrays, but you don't really need it for that.
- Returned memory is filled with zeros
 - So, calloc() may be a little more expensive.

Performance Comparison

- We have three storage regions we can use
 - each with different behavior
 - .. and performance characteristics
- Let's try out some different techniques and see what they cost.

Uninitialized Stack Storage

```
void f() {  
    int a[ 1000 ];  
    a[ 999 ] = 1;  
    counter += a[ 999 ];  
}  
  
int main( void )  
{  
    for ( long i = 0; i < 100000000; i++ )  
        f();  
    return 0;  
}
```

Make a big array on the stack,
but don't initialize it.

Just use it a little.

Using this to modify a global
variable (so the compiler
doesn't optimize-out all this
code)

Do this 100,000,000 times.

Runtime: 0.42 seconds

Initialized Stack Storage

```
void f() {  
    int a[ 1000 ] = {};  
    a[ 999 ] = 1;  
    counter += a[ 999 ];  
}  
  
int main( void )  
{  
    for ( long i = 0; i < 100000000; i++ )  
        f();  
    return 0;  
}
```

Same thing, but initializing
the array.

Runtime: 9.22 seconds

Uninitialized stack space is cheap,
but initializing it has a cost.

Uninitialized Heap Storage

```
void f() {  
    int *a = malloc( 1000 * sizeof( int ) );  
    a[ 999 ] = 1;  
    counter += a[ 999 ];  
    free( a );  
}  
  
int main( void )  
{  
    for ( long i = 0; i < 100000000; i++ )  
        f();  
    return 0;  
}
```

This time, on the heap,
uninitialized.

Runtime: 6.31 seconds

Allocating stack space is much
cheaper than heap.

Initialized Heap Storage

```
void f() {  
    int *a = calloc( 1000, sizeof( int ) );  
    a[ 999 ] = 1;  
    counter += a[ 999 ];  
    free( a );  
}  
  
int main( void )  
{  
    for ( long i = 0; i < 100000000; i++ )  
        f();  
    return 0;  
}
```

On the heap, initialized this time.

Runtime: 16.17 seconds

Initializing any memory has a cost.

Static Storage

```
int a[ 1000 ];
```

Created once at the start of execution.

```
void f() {  
    a[ 999 ] = 1;  
    counter += a[ 999 ];  
}
```

Probably, this is what makes it faster.

```
int main( void )  
{  
    for ( long i = 0; i < 100000000; i++ )  
        f();  
    return 0;  
}
```

Runtime: 0.28 seconds

You may be able to compute static addresses at compile time (depending on the OS and hardware)

Now in Java

```
public class Test {  
    static long counter = 0;  
  
    static void f() {  
        int[] a = new int [ 1000 ];  
        a[ 999 ] = 1;  
        counter += a[ 999 ];  
    }  
  
    public static void main( String[] args ) {  
        for ( long i = 0; i < 100000000; i++ )  
            f();  
    }  
}
```

In Java, all arrays are on the heap

... and they're initialized to zero

... and they're eventually garbage collected.

Runtime: 97.1 seconds

This all has a cost.