

# Bit Operations

CSC 230 : C and Software Tools

NC State Department of Computer  
Science

# Topics for Today

- Leftover topic: Linked lists with pointer-to-pointer
- Bitwise operator review
- Bitwise tasks
  - Setting / clearing / testing selected bits
  - Copying selected bits
  - Movable masks
  - Manipulating a bit field
- Bit fields in structs

# Typedef and 2D Arrays

- This can also simplify describing multi-dimensional arrays.

```
typedef int Row[ 20 ];
```

```
Row *table;
```

```
table = (Row *) malloc( 50 * sizeof( Row ) );
```

I'm a pointer to a Row, so you can use me like an array of Rows.

Now, I'm an array of 50 rows.

- We could do without this typedef, but it makes the types harder to describe.

```
int (*table)[ 20 ];
```

```
table = (int (*)(20)) malloc( 50 * 20 * sizeof( int ) );
```

Oops. This is just the same thing without typedef.

# Linked List Representation

- Same representation as before.

```
struct NodeTag {  
    int value;  
    struct NodeTag *next;  
};  
  
typedef struct NodeTag Node;
```

Representation for a  
Node.

```
typedef struct {  
    Node *head;  
} List;
```

Representation for a  
whole list.


# Removing Nodes, Classic

- Normally, to remove a node we might handle the head of the list as a special case:


```
bool removeValue( List *list, int val )
{
    if ( list->head && list->head->value == val ) {
        Node *n = list->head;
        list->head = list->head->next;

        free( n );
        return true;
    }
    ...
}
```

First node as a special case.



# Removing Nodes, Classic



```
...;  
Node *pred = list->head;  
while ( pred->next && pred->next->value != val )  
    pred = pred->next;  
  
if ( pred->next ) {  
    Node *n = pred->next;  
    pred->next = pred->next->next;  
  
    free( n );  
    return true;  
}  
  
return false;  
}
```

Look for the node  
**before** the one you  
want to remove.

Unlink (and free) it.

# An Idea that Almost Works

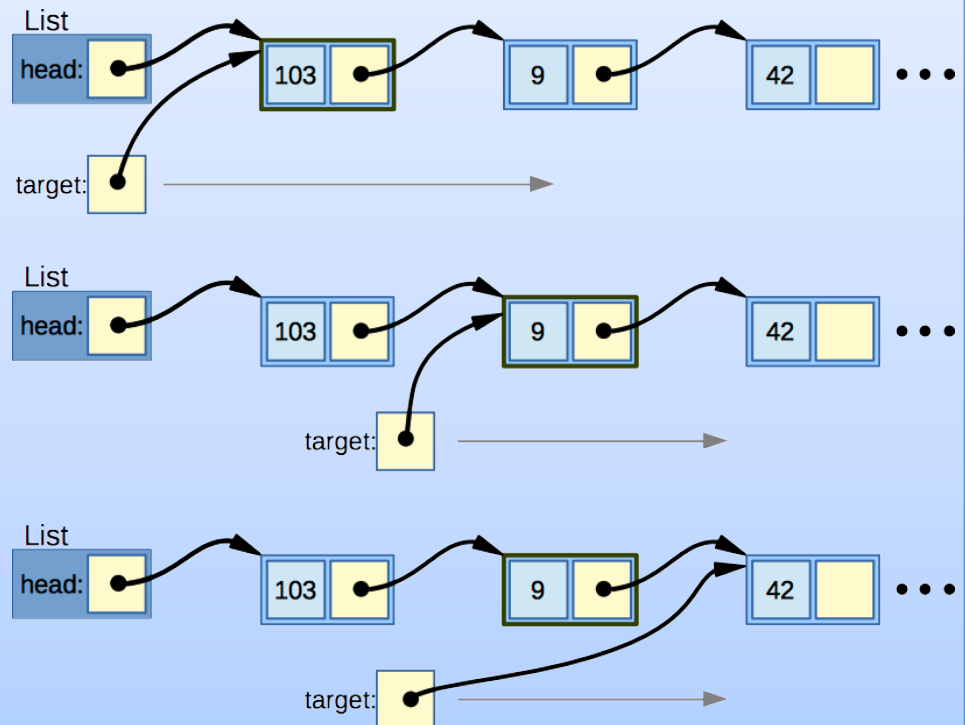
- How about this technique?

```
bool removeValue( List *list, int val )
{
    Node *target = list->head;
    while (target &&
           target->value != val)
        target = target->next;

    if ( target ) {
        Node *n = target;
        target = target->next;


        free( n );
        return true;
    }

    return false;
}
```



# Removing Nodes, “Simplified”

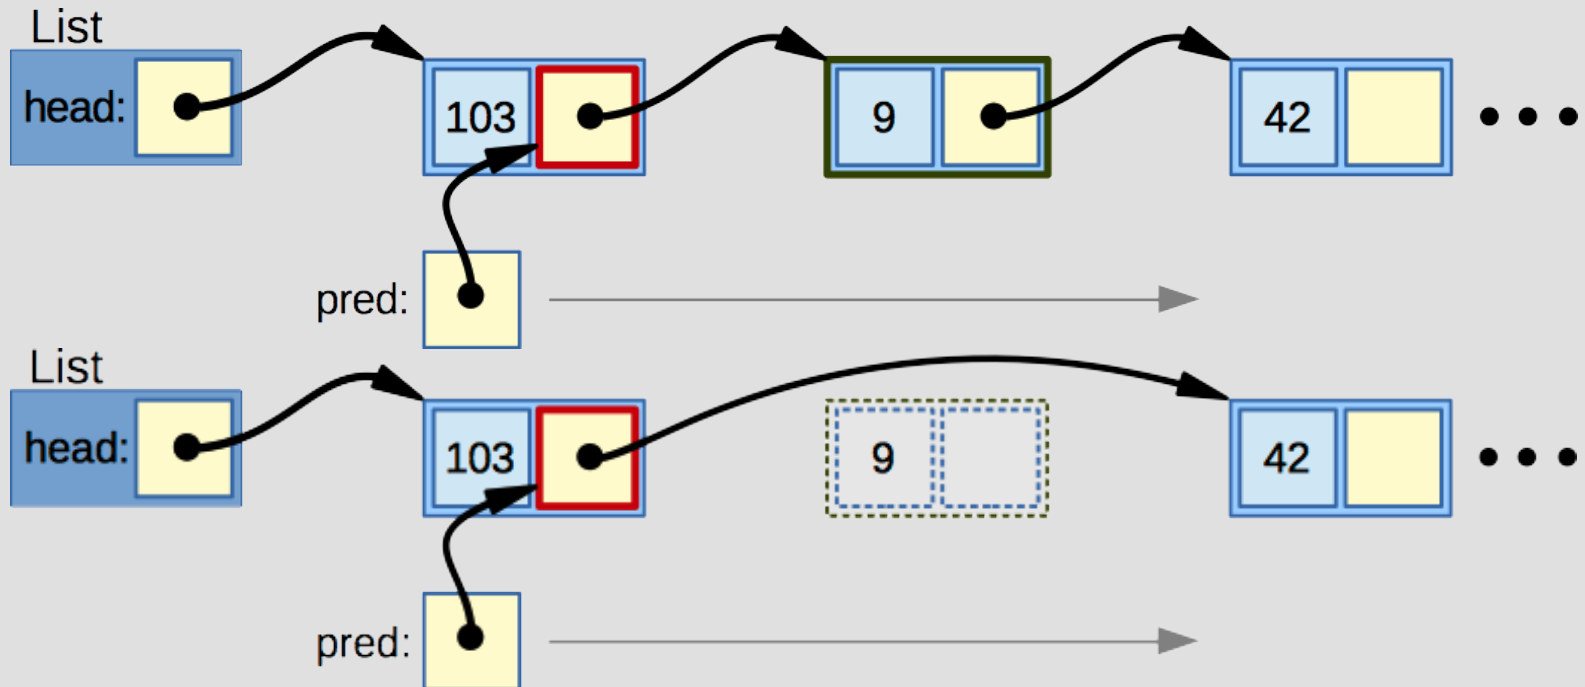
- To remove a node, we must change the pointer that points to it.
  - Every node has such a pointer
  - Either **the pointer inside its predecessor**.
  - **Or, the head pointer** itself.
- Both of these pointer have the same type (pointer to Node)
- We can handle them uniformly via a pointer to pointer to Node.
- We’ll start with a pointer pointing to the head pointer:





# Removing Nodes, “Simplified”

- As we traverse the list, we will move the ahead to pointers within the nodes.
  - This will give us a way to remove the next node on the list.



# Removing Nodes, Simplified

```
bool removeValue( List *list, int val )
{
    Node **pred = &list->head;
    while ( *pred && (*pred)->value != val )
        pred = &(*pred)->next;

    if ( *pred ) {
        Node *n = (*pred);
        *pred = (*pred)->next;


        free( n );
        return true;
    }

    return false;
}
```

I'm the address of the pointer to the node you're thinking about removing.

Start at the head, then walk down through all the pointers inside the nodes.

If we find a node containing val, remove its node by changing the pointer that points to it.



# Looking Inside a Value

- Internally, int and double values are represented with patterns
- A good programming language lets us ignore details of how this is done
  - We can think of an int as ... well, an integer
- But, sometimes we want to work with individual bits in a value ... why?
  - Maybe we could pack more values into limited memory.
  - Maybe we need to talk to hardware or work with a file format that requires particular bit patterns
  - Maybe we need to perform compression or encryption
  - Maybe we can better exploit bit-level parallelism in the hardware

# Working with Bits

- C99 provides **no standard way** to ...

- ... give a literal constant in binary

```
i = 011001011;
```

- ... read an integer as a string of binary digits

```
scanf( "%b", &i );
```

- ... print a value as a string of binary digits

```
printf( "%b", i );
```

- What can we do?

- Use octal or hexadecimal
- Write our own code to work at the bit level.

# C Bit-Level Operators

- We have some operators just for manipulating values at the bit level

Operator	Description
++ --	Postincrement, postdecrement
++ --	Preincrement, predecrement
+ -	Unary positive, negative
!	Logical not
~	Bitwise complement
(type)	Type cast
sizeof	Well, it's the sizeof operator
* / %	Multiply, divide, mod
+ -	Add, subtract
<< >>	Bitwise left and right shift
< <= > >=	Relational operators

# More C Operators

Operator	Description
== !=	Equals, not equals
&	Bitwise and
^	Bitwise exclusive or
	Bitwise (inclusive) or
&&	Logical and
	Logical or
? :	Ternary operator
= += -= *= ...	Assignment, and its friends
,	Comma operator

# Working With Bits

- We're about to start working with the binary representation of values
- We're going to notice some behaviors we normally don't have to worry about.
  - This one is called sign extension.

```
unsigned char x;  
signed char y;
```

```
x = 0xA9; // binary 10101001 (that means 169)  
y = 0xA9; // binary 10101001 (here, that's -87)
```

```
printf( "%x\n", x );  
printf( "%x\n", y );
```

This prints a9.  
Good

But, this prints  
ffffffa9

But, that's  
how int -87 is  
represented.

# Meet ~ (a Unary Bitwise Operator)

- Unary operator `~` is a **bitwise complement**
- It's like `-` or `!`, but it works **across all bits**

```
unsigned char x, y, z;
```

```
x = 0x1A; // binary 00011010
```

```
y = ~x;    // binary 11100101
```

[illegible]

- See,  $\sim$  is different from these other operators

```
signed char x, y, z, w;
```

```
x = 0x1A; // binary 00011010
```

```
y = ~x;    // binary 11100101
```

```
z = !x;    // binary 00000000
```

```
w = -x; // binary 11100110
```

Same here, it's  
1111111111111111111111111100110  
until you assign it.



# Binary Bitwise Operators

- The next three operators are binary, they take two operands
  - **&** is a bitwise **and**  
a 1 in the result only if both corresponding input bits 1
  - **^** is bitwise **exclusive or**  
a 1 in the result only if both corresponding input bits are different
  - **|** is bitwise **(inclusive) or**  
a 1 in the result if either of the corresponding input bits are 1

# Binary Bitwise Operators

- Here's how they work:

```
  0 1 1 0 1 1 0 0
& 1 1 0 0 1 0 1 0
-----
  0 1 0 0 1 0 0 0
```

```
  0 1 1 0 1 1 0 0
^ 1 1 0 0 1 0 1 0
-----
  1 0 1 0 0 1 1 0
```

```
  0 1 1 0 1 1 0 0
| 1 1 0 0 1 0 1 0
-----
  1 1 1 0 1 1 1 0
```

- See, these are different from our old friends.

```
  0 1 1 0 1 1 0 0
&& 1 1 0 0 1 0 1 0
-----
  0 0 0 0 0 0 0 1
```


I'm drawing these like arithmetic problems from elementary school. It's easier to see the bits if they're lined up.

```
  0 1 1 0 1 1 0 0
|| 1 1 0 0 1 0 1 0
-----
  0 0 0 0 0 0 0 1
```

# Shift Operators

- Left logical shift of  $x$  by  $y$  bits:  $x \ll y$ 
  - For example:  $0x1A \ll 5$

If you shift far enough, you can lose bits off the high end.




New zeros shifted in on this end.

- Operands should be integer types
- If  $x$  has  $n$  bits,  $y$  should be between 0 and  $n-1$
- You can't left shift by a negative amount, but ...

# Shift Operators

- **Right logical** shift of **x** by **y** bits: **x** >> **y**
  - For example: 0xCA >> 2

You get new high-order zeros shifted in on the left.





bits shifted past the low-order bit are discarded.

- Again, operands should be integer types
- If **x** has **n** bits, **y** should be between 0 and **n**-1
- Really, what gets shifted in on the depends ....

# Shift Operators

- Right **arithmetic** shift of **x** by **y** bits: **x** **>>** **y**
  - Wait! This is the same operator.
  - Behavior depends on whether the left operand is **signed**.

For signed operands, you get copies of the high-order bit shifted in.



# Shift Operators

- Why arithmetic shift?
  - For signed values, this maintains the sign.
- We can use left shift to multiply by 2 or any power of 2
  - $x \ll y$  is the same as  $x * 2^y$
- And, we can use right shift like a divide operation
  - $x \gg y$  is the same as  $x / 2^y$

```
unsigned int x y;
```

```
x = 75 << 2; // that's 300
```

```
y = 75 >> 2; // that's 18
```

But, this only works for positive numbers.

# OK, How Do We Use These Operators?

- What can we do with these operators?
- Clear selected bits to zero or set them to 1
- Test if selected bits are zero or 1
- Print or count 0 and 1 bits
- Copy selected bits from x to y or to another position in x
- Access a bit field as if it was a tiny integer.

# Clearing Selected Bits

- We can think of the & operator as applying a *mask*

	0	1	1	0	1	1	0	0
&	1	1	0	0	1	0	1	0
-----								
	0	1	0	0	1	0	0	0

I'm the value you want to work with.

I'm the mask, deciding what bits to keep and discard.

- Wherever we have a 0 in the mask, a bit is cleared
- Wherever we have a 1, the original value passes through



# Clearing Selected Bits

- So, we could clear just the low-order 4 bits

```
unsigned char x;  
x = 0xD3;           // binary 11010011  
x = x & 0xF0;       // mask   11110000  
                    // now    11010000
```


- Or, we could clear every other bit.

```
x = 0xD3;           // binary 11010011  
x &= 0x55;          // mask   01010101  
                    // now    01010001
```

We also have these for  
bitwise operators.

# Setting Selected Bits

- Again, it's just how you think about the bitwise operators
- We can think of `|` as applying a different kind of mask



- Wherever we have a 1, a bit is set
- Wherever we have a 0, the original value passes through

# Setting Selected Bits

- So, we could set just the low-order 4 bits

```
unsigned char x;  
x = 0xD3;           // binary 11010011  
x = x | 0x0F;       // mask   00001111  
                    // now    11011111
```

- Or, we could set just the first and last.

```
x = 0xD6;           // binary 11010110  
x |= 0x81;          // mask   10000001  
                    // now    11010111
```

# Inverting Selected Bits

- With the ^ operator, a 1 in the mask says where to flip a bit

0	1	1	0	1	1	0	0
^	1	1	0	0	1	0	1
-----							
1	0	1	0	0	1	1	0

I'm the value you want to work with.

I say what bits to change.

- We could use this to flip just the low-order bit

```
unsigned char x;  
x = 0xD3;           // binary 11010011  
x ^= 0x01;          // mask   00000001  
                     // now    11010010
```

# Testing Selected Bits

- First, use & to choose the bits you want.

	0	1	1	0	1	1	0	0
&	1	1	0	0	1	0	1	0
-----								
	0	1	0	0	1	0	0	0

- Then what?
  - If the result is zero, none of those bits were 1
  - If the result matches the mask, they were all 1

Careful here,  
precedence.

```
unsigned char x, mask;  
x = 0xD3;           // binary 11010011  
mask = 0x0F;        // mask   00001111  
if ( ( x & mask ) != 0 ) // result 00000011  
    ...;  
if ( ( x & mask ) == mask ) // result 00000011  
    ...;
```

# Counting or Printing Bits

- So, if we can build a mask, we can get to any bits we want.
- What if we want to be able to access every bit ... individually.
  - Say, we want to iterate over the bits in an int.
- We can build a **movable mask** for this
  - $0x01 \ll i$  is a mask with a 1 just in bit position  $i$
  - And, of course,  $\sim(0x01 \ll i)$  is the complementary mask

Is this the same as  
 $\sim 0x01 \ll i$  ?

# Counting Bits

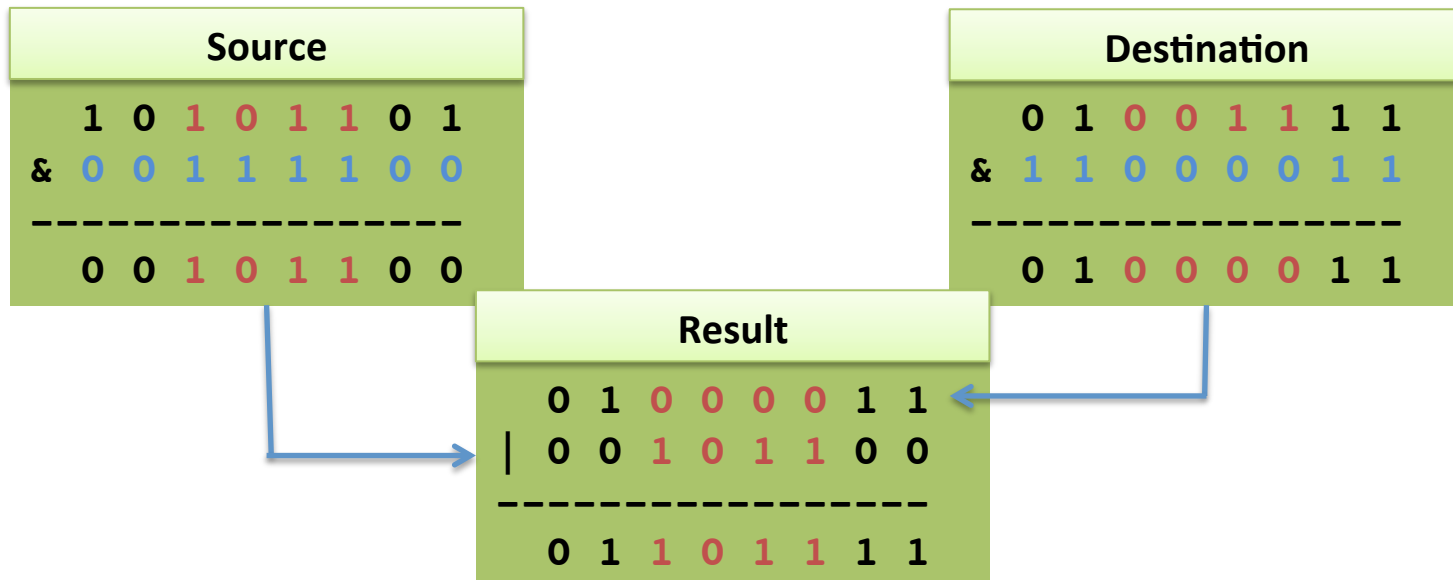
- We can iterate over the bits, then apply the mask to check each one.

```
unsigned int bill = 2348291;
int bcount = 0;
for ( int i = 0; i < 8 * sizeof( int ); i++ ) {
    if ( bill & 1 << i )
        bcount++;
}
printf( "That's %d one bits\n", bcount );
```

- We could use the same technique to print a number in binary.

# Copying Selected Bits


- Say we wanted to copy the middle four bits in a character
- In the source, we can use & to mask off everything but the bits we want to copy.
- In the destination, we can use & with a complementary mask to clear out these bits
- Then, we can use | to turn on just the bits copied from the source






# Bit Fields

- We can pack multiple small integer values into a variable.
  - This can save some storage or reduce communication overhead.
  - It can even let us manipulate multiple values in parallel.



# Extracting a Bit Field

- To extract a particular field:
  - Mask off the other bits
  - Then, shift down (so you can use the contents of the field like an int)
  - Say, we want to get just the middle 4 bits of a char.



# Changing a Bit Field

- Then, we can manipulate the value of that field.
- And put the result back when we're done.

	1	0	1	0	1	1	0	1
&	1	1	0	0	0	0	1	1
-----								
	1	0	0	0	0	0	0	1

Original char,  
make room for  
the new field  
value.

Insert the  
updated value

	1	0	0	0	0	0	0	1
	0	0	1	1	0	0	0	0
-----								
	0	1	1	1	0	0	0	1

			1	0	1	1	
+							1
-----							
			1	1	0	0	
<<							2
-----							
			1	1	0	0	0
&	0	0	1	1	1	1	0
-----							
	0	0	1	1	0	0	0

Add one to  
the field.

Shift it back  
to where it  
goes.


Mask off other  
bits (in case of  
overflow)

# Bit Fields in Structs

- Inside a struct, we can use integer fields with variable numbers of bits
  - The compiler will worry about packing/unpacking them in words of memory.
- Why would we want this?
  - We could save memory, if we need several fields for small integer values
  - Maybe we could match the binary format of some file
  - ... or the format for communicating with some hardware device


# Using Bit Fields

- Normally, fields occupy consecutive groups of whole bytes, maybe even with some padding.
- Internally, bit fields can use parts of individual bytes.



```
typedef struct {  
    unsigned short red;  
    unsigned short green;  
    unsigned char blue;  
    unsigned char alpha;  
} RegularColor;
```

The compiler says "6 Bytes"



```
typedef struct {  
    unsigned int red:9;  
    unsigned int green:9;  
    unsigned int blue:6;  
    unsigned int alpha:6;  
} PackedColor;
```

But here just 4.

Here's the  
new syntax

# Using Bit Fields

- Mostly, you can use bit fields like any other integer field.

```
PackedColor c1 = { 511, 256, 0, 32 };
```

Initialize them.

```
int r = c1.green;
```

Read them.

```
c1.blue += 16;
```

Assign and change them.

```
c1.red++;
```

Overflow them.

- But, you can't take their addresses. A bit field may start inside a byte.

```
void *ptr = &( c1.red );
```



# Wait Just a Minute!

- Didn't we just learn how to do bit fields, with bitwise operators!
- Was that a waste of time?
  - ~~Yes~~ ... I mean No.
- For bit fields in a struct, the compiler determines the layout.
  - So code that depended on a particular packing of the bits wouldn't be portable.
  - And if we really wanted to use a file format or talk to hardware with bit fields in a particular format ...
  - ... we'd probably have to write that code ourselves

# Wait Just a Minute!

- So, are bit fields in a struct useful?
- Yes
  - They can reduce the memory footprint of your program.
  - They can give you convenient access to bit fields in a file, if portability isn't important.