

More Standard Library and Program Organization

CSC 230 : C and Software Tools
NC State Department of Computer
Science

Topics to Discuss

- Environment Variables
- Variable numbers of parameters
- Nonlocal jumps
- Program structure
 - Abstract Data Types
 - Generic ADTs

Environment Variables

- Your execution environment supports variables
 - We can access these from the shell

```
$ echo $HOME  
/home/jerry  
$ NAME=linda  
$ echo $NAME  
linda
```

- The shell uses these a lot for configuration and to customize behavior
 - HOME, PATH, SHELL, PS1, EDITOR, ...
- But, these aren't just available to the shell, all programs have environment variables.

Using the Environment

- The standard library lets us get/set these variables
 - We need our good friend, `stdlib.h`
 - Then, we can gen environment variable values

```
char *str = getenv( "HOME" );
printf( "HOME = %s\n", str );
```

- And, we can change them

```
setenv( "HOME", "/some/other/place", 1 );
```

Variable to set.

new value.

Should I overwrite?

- These changes are local to the program, they go away when it exits.

Variable Numbers of Arguments

- We can define functions that take a variable number of arguments

- printf() needs this
 - The notation looks like this:

```
int printf( const char *fmt, ... )
```

“I take any number of additional arguments.”

- Typically, you'll have:

- One or more **named arguments** (at least one is required)
 - Followed by a variable number of **unnamed arguments**

Processing Variable Arguments

- We can handle variable numbers of arguments ... but we'll need some help.

```
#include <stdarg.h>
```

- We'll need to do four things
 - Create a variable of type `va_list`, to store vararg state
 - Call `va_start()`, to start processing variable arguments
 - Call `va_arg()` for each argument, to extract it with the right type
 - Call `va_end()` when we're done

Counting Arguments

- There's no automatic mechanism to tell how many arguments we got
- We'll need to build-in a way to do that
 - Like printf(), we could use one of the named arguments to figure it out
 - Or, we could use a mechanism to mark the last argument
- Let's try both of these ...

Adding Multiple Arguments

```
int addIntegers( int n, ... )  
{  
    va_list ap; _____| State for vararg processing.  
    va_start( ap, n ); _____| Start processing after n.  
    int total = 0;  
    for ( int i = 0; i < n; i++ )  
        total += va_arg( ap, int );  
    va_end( ap ); _____| Get the next int argument.  
    return total;  
}
```

```
int sum = addIntegers( 5, 17, 42, 98, 3, 27 );
```

```
int sum2 = addIntegers( 29, 17, 42, 98, 3, 27 );
```

It's our job to use the function properly.

Printing Multiple Strings

```
void printStrings( char *first, ... )  
{  
    va_list ap;  
    va_start( ap, first );  
  
    printf( "%s\n", first );  
  
    char *str;  
    while ( ( str = va_arg( ap, char * ) ) ) {  
        printf( "%s\n", str );  
    }  
  
    va_end( ap );  
}
```

We must have at least one named parameter.

Keep going until we hit the NULL.

```
printStrings( "one", "two", "three", "four", NULL );
```

A null to mark the end.

Max of Multiple Values

```
double maximum( int n, double v1, ... )  
{  
    va_list ap;  
    va_start( ap, v1 );  
  
    double largest = v1;  
  
    for ( int i = 1; i < n; i++ ) {  
        double v = va_arg( ap, double );  
        if ( v > largest )  
            largest = v;  
    }  
  
    va_end( ap );  
    return largest;  
}
```

How many values?

Force the caller to give us at least one.

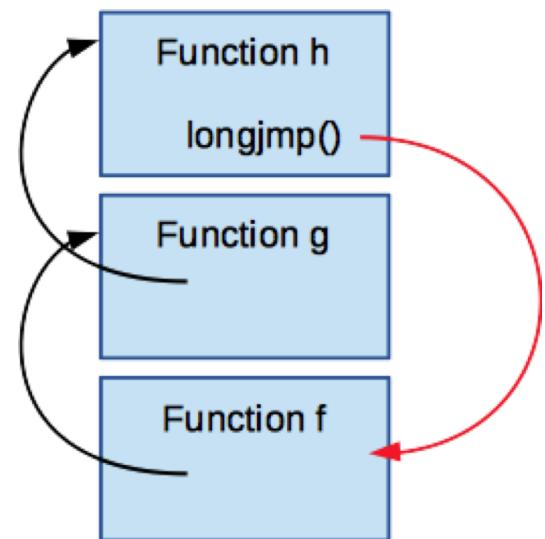
Handle the first value as a named argument.

Handle the rest through varargs.

```
max = maximum( 6, 32.1, 15.32, 2.99, 45.2, 13.62, 7.0 );
```

Nonlocal Jumps

- A goto will let you jump anywhere
 - ... as long as it's in the same function.
 - Maybe an OK mechanism for handling errors
- C also supports a long jump
 - To jump out of a function back to one of its callers
 - Made for handling errors
 - Kind of like an exception in Java



Nonlocal Jumps

- Supported by the `setjmp.h` header.
- Before you can `longjmp`, you must take a picture of the place you want to jump back to.
 - This is called your program's `environment`
 - Where it's executing
 - And where the top of the stack is
- The environment must be stored in instance of `jmp_buf`

```
jmp_buf env;
```

Taking a Picture

- The setjmp() macro saves the state of the environment.

```
setjmp( env );
```

- Initially, it returns zero.
- You need to check this (for later).

```
if ( setjmp( env ) == 0 ) {  
    ..
```

Hitting Rewind

- If something goes wrong, calling `longjmp()` will take you back to that picture.

```
if ( something_went_wrong )
    longjmp( env, 1 );
```

Some int you can use to indicate what went wrong.

- You can allocate the environment statically
- Or pass it to the function for it to use.
- But, you can only `longjmp()` up the stack.

Waking up in setjmp()

- The longjmp() takes you right back to the matching setjmp() call

```
if ( something_went_wrong )
    longjmp( env, 1 );
```

Some other function.

```
jmp_buf env;
if ( setjmp( env ) == 0 ) {
    Call some other function.
```

- So it looks (almost) like you're right back where you started.

Waking up in setjmp()

- After a longjmp(), the setjmp() won't return zero.
 - It returns the value you passed to longjmp()

```
if ( something_went_wrong )
    longjmp( env, 1 );
```

Some other function.

```
jmp_buf env;
if ( setjmp( env ) == 0 ) {
    Call some other function.
} else {
    Error handling code.
}
```

The diagram illustrates the control flow between two functions. A blue oval encloses the first code block (the longjmp call). An arrow points from this oval to the first line of the second code block (the setjmp check). Another arrow points from the 'Call some other function.' line to the 'Error handling code.' line, indicating that if the setjmp check fails (returns 0), the program will fall through to the error handling code.

I'll return 1 this time.

Using setjmp()/longjmp()

- We can use this kind of like an exception.

```
void middleFunction( jmp_buf *eenv )
{
    int sum = readValues( eenv );
    printf( "Sum: %d\n", sum );
}
```

Plan ahead for where to go if something fails.

```
void topFunction()
{
    jmp_buf eenv;

    if ( setjmp( eenv ) == 0 ) {
        middleFunction( &eenv );
    } else {
        printf( "Invalid input!\n" );
    }
}
```

Passing the error-handling environment among functions.

An error in here could jump past this middle function.

Using setjmp()/longjmp()

- Here's where we notice the error.

```
int readValues( jmp_buf *eenv )
{
    int sum = 0;
    int m, v;

    while ( ( m = scanf( "%d", &v ) ) == 1 )
        sum += v;

    if ( m != EOF )
        longjmp( *eenv, 1 );

    return sum;
}
```

Jump back to the error-handling code.

Skipping the printout in the middle function.

Cleaning Up

- What if we had some resources to free?

```
int readValues( int *list, int capacity, FILE *fp, jmp_buf *eenv )
```

```
void middleFunction( jmp_buf *eenv )
{
    FILE *fp = fopen( "input.txt", "r" );
    if ( !fp )
        longjmp( *eenv, 1 );

    int capacity = 100;
    int *list = (int *)malloc( capacity * sizeof( int ) );

    int vcount = readValues( list, capacity, fp, eenv );

    for ( int i = vcount - 1; i >= 0; i-- )
        printf( "%d\n", list[ i ] );

    fclose( fp );
    free( list );
}
```

Now, this function reads from a file ... into an array.

A jump past this function would leak these resources.

We need something like a finally block for this.

Cleaning Up

- We can catch and re-throw to free resources.
- With no changes to the other functions.

```
void middleFunction( jmp_buf *parentEnv )
{
    FILE *fp = fopen( "input.txt", "r" );
    if ( !fp ) longjmp( *parentEnv, 1 );
    int *list = (int *)malloc( 100 * sizeof( int ) );

    jmp_buf midEnv;
    if ( setjmp( midEnv ) == 0 ) {
        int vcount = readValues( list, capacity, fp, &midEnv );
        ... print the list ...
        fclose( fp );
        free( list );
    } else {
        fclose( fp );
        free( list );
        longjmp( *parentEnv, 1 );
    }
}
```

Our own environment, to catch the jump.

Tell the readValues() function to go here on an error.

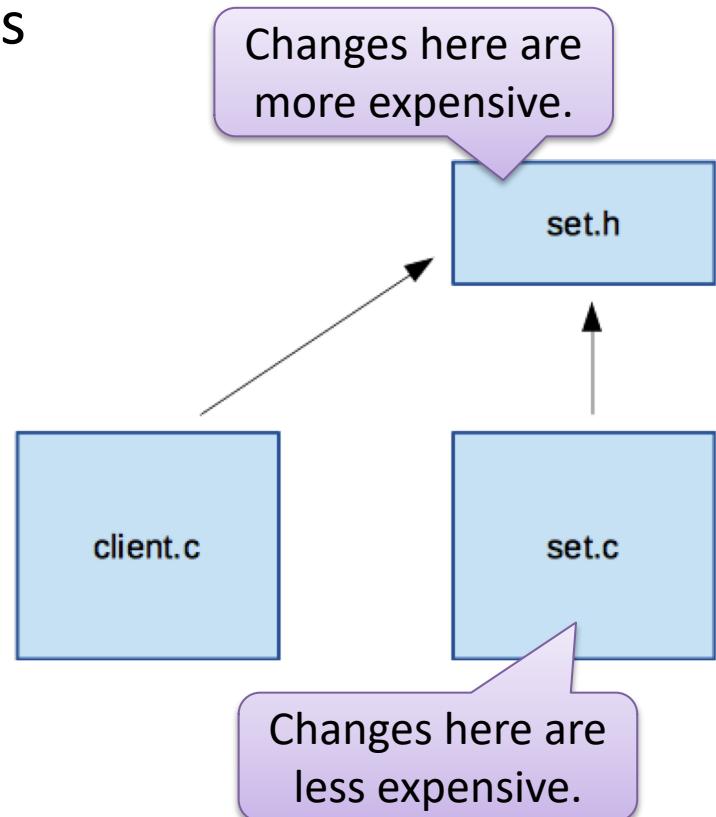
Clean up our part, and jump on up the stack.

Designing a Component

- We're going to look at a few alternatives for designing and implementing a component.
- It will support a set of (integer) values.
 - makeSet() : to initialize the representation
 - insertValue() : to add a value to the set
 - containsValue() : to report whether a value is in the set.
 - destroySet() : to free any resources used by the set.

Designing a Component

- The set component will have a header to advertise its interface.
- And an implementation file (to implement that interface)
- *Client code* is anything else that uses this component.
- We'd like to put as little as possible in the header
 - Why? If something's in the header, we can't change it without recompiling (all) the client code.
 - Imagine if someone had to fix a bug in stdio.h!



Hiding the Implementation

```
void makeSet();  
void insertValue( int val );  
bool containsValue( int val );  
void destroySet();
```

set.h

```
#include "set.h"  
  
struct Node {  
    int val;  
    struct Node *next;  
};  
static struct Node *head;  
  
void makeSet()  
{  
    head = NULL;  
}  
  
void insertValue( int val )  
{  
    for ( struct Node *n = head; *n; n = n->next )  
        ...
```

set.c

Hidden Representation

We could change this representation without breaking client code.

Client Code

```
#include "set.h"

int main()
{
    makeSet();

    insertValue( 2 );
    insertValue( 3 );
    insertValue( 5 );
    insertValue( 7 );
    insertValue( 11 );

    for ( int i = 0; i <= 12; i++ )
        if ( containsValue( i ) )
            ...

    destroySet();
}
```

Initialize the set

Perform some operations
with it.

Clean up when we're
done.

Abstract Object

- The textbook would call this an abstract object.
 - You can use this object without being sensitive to how it's implemented.
- But, here, we can only have one set.
 - We'd like client code to be able to make as many sets as they want.
 - We need to give client code a type that represents a set.

Exposing the Implementation

```
struct Node {  
    int val;  
    struct Node *next;  
};  
  
typedef struct {  
    struct Node *head;  
} Set;  
  
void makeSet( Set *s );  
void insertValue( Set *s, int val );  
bool containsValue( Set *s, int val );  
void destroySet( Set *s );
```

set.h

Here, use this type for each set you need.

Every function takes a (pointer to the) set it's supposed to operate on.

```
#include "set.h"
```

set.c

```
void makeSet( Set *s )  
{  
    s->head = NULL;  
}  
  
...
```

Client Code

```
#include "set.h"

int main()
{
    Set s;
    makeSet( &s );

    insertValue( &s, 2 );
    insertValue( &s, 3 );
    insertValue( &s, 5 );
    insertValue( &s, 7 );
    insertValue( &s, 11 );

    for ( int i = 0; i <= 12; i++ )
        if ( containsValue( &s, i ) )
            ...

    destroySet( &s );
}
```

Make an instance of set.

Initialize it, passing the address of the uninitialized set we just made.

Perform some operations with it.

Clean up when we're done.

A Bad Idea

```
#include "set.h"

int main()
{
    Set *s;
    makeSet( s );

    insertValue( s, 2 );
    insertValue( s, 3 );
    insertValue( s, 5 );
    insertValue( s, 7 );
    insertValue( s, 11 );

    for ( int i = 0; i <= 12; i++ )
        if ( containsValue( s, i ) )
            ...

    destroySet( s );
}
```

All these functions want a pointer to a Set, so that's what I'll give them.

This will compile ... but I bet it will crash when you run it.

No More Abstraction?

- To let client code make instances of Set, we had to let it see our representation.
- This is bad
 - Maybe our design is super secret (less important)
 - Maybe this makes the header larger, which will slow compilation for all code using it (less important)
 - Maybe we want to be able to change our design or fix our implementation later (more important)
- Can we let client code use the Set type ... while keeping its contents secret?

Incomplete Type

- Yes. We can give what's called a *incomplete type*.
 - We can tell the compiler about a struct, without saying what's in it.

```
struct ListStruct;
```

There's a structure named this.

```
typedef struct RectangleStruct Rect;
```

There's a structure named this,
and here's a type name for it.

- This permits name-only use of the struct.
 - You can work with pointers to it.
 - ... but you can't do anything that depends on its size or its contents.

Re-Hiding the Implementation

- We can use an incomplete type to insulate client code from the implementation.
- But, client code can only use this type as an *opaque pointer*.
- We'd call this design an *Abstract Data Type* (ADT).
- Now, the header might look like:

```
typedef struct SetStruct Set;  
  
Set *makeSet();  
void insertValue( Set *s, int val );  
bool containsValue( Set *s, int val );  
void destroySet( Set *s );
```

Here, use this type when you need a Set.

But, I'm not telling you anything about it. You have to ask me to make one for you.

Looking Inside

```
#include "set.h"

struct Node {
    int val;
    struct Node *next;
};

struct SetStruct {
    struct Node *head;
};

Set *makeSet()
{
    Set *s = (Set *) malloc( sizeof( Set ) ) ;
    s->head = NULL;
    return s;
}
...
void destroySet( Set *s )
{
    ...
    free( s );
}
```

set.c

In the implementation file, we provide the type we provide the full type definitions.

I'll have to give you a dynamically allocated instance (so, more expensive than stack allocation).

And, you'll need my help when you're done with this object.

Client Code

```
#include "set.h"

int main()
{
    Set *s = makeSet();

    insertValue( s, 2 );
    insertValue( s, 3 );
    insertValue( s, 5 );
    insertValue( s, 7 );
    insertValue( s, 11 );

    for ( int i = 0; i <= 12; i++ )
        if ( containsValue( s, i ) )

    destroySet( s );
}
```

Please make one of those Set objects I've heard about.

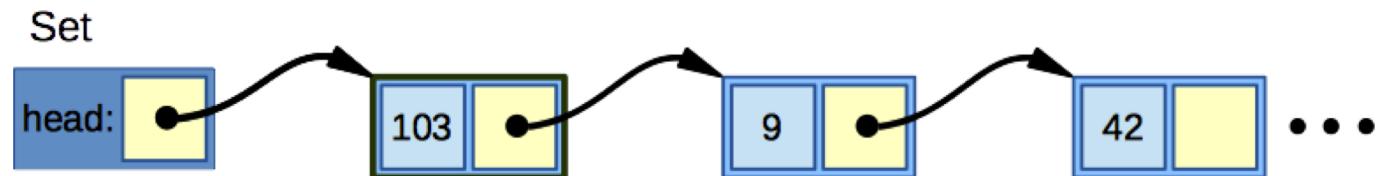
The rest of this is the same, except s is now a pointer.

A Generic, Abstract Data Type

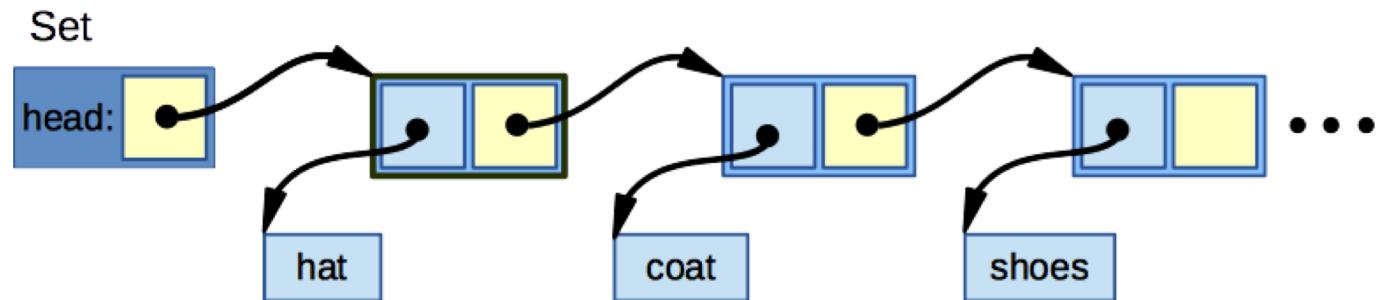
- All we have is a set of ints
- What if we wanted a set of something else?
 - We'd have to make another component
 - ... but much of its code would be the same
 - If we were clever, we could make this code generic.
- We would call this a *Generic ADT*

A Generic, Abstract Data Type

- In the implementation, we'd have to be able to store any type of data in our set.
- Previously, we stored a fixed-sized value right inside each node.

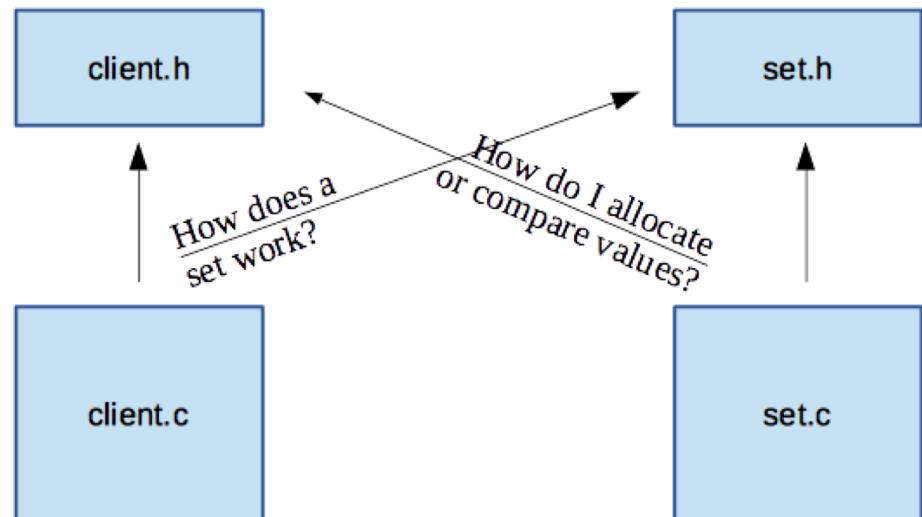


- Now, we'll use a generic pointer to each value, stored elsewhere in memory.



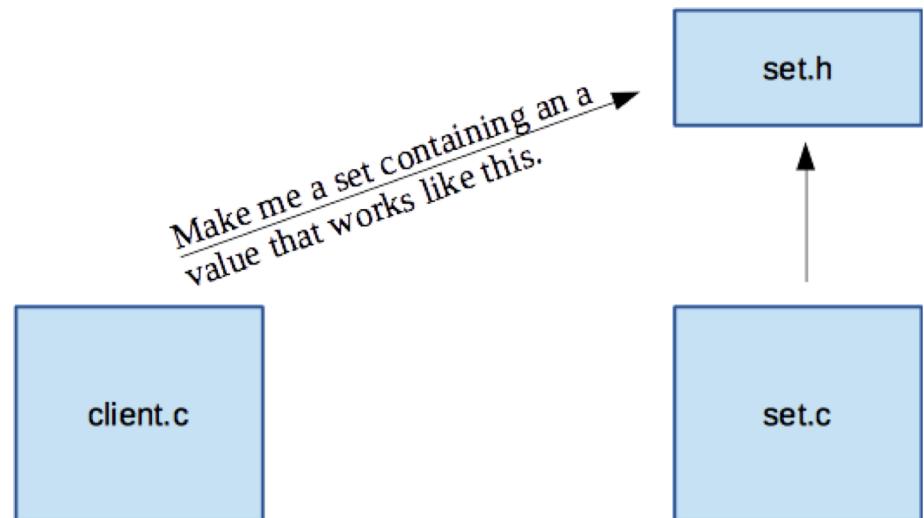
Circular Dependency

- How will our set know
 - How what value type it's supposed to contain?
 - Or how to compare values of that type?
- Could it just use code from the client for this?
- No. This is a terrible idea.
- It creates a *circular dependency* in our code.



Avoiding Circular Dependency

- Lots of design cleverness involves eliminating circular dependency.
- Our set needs to provide a mechanism for telling it about the value it's working with.
- Kind of like what qsort()
does
- ... it provides a way to
tell it something about
the type of values it's
sorting.



Generic Implementation

```
#include "set.h"

struct Node {
    void *val;
    struct Node *next;
};

struct SetStruct {
    size_t vsize;
    bool (*same)(void const *v1, void const *v2);
    struct Node *head;
};

Set *makeSet(size_t vsize, bool (*same)(void const *v1, void const *v2))
{
    Set *s = (Set *) malloc( sizeof( Set ) );
    s->head = NULL;

    s->vsize = vsize;
    s->same = same;

    return s;
}
```

set.c

I could hold any value.

I have to remember how large values are, and how to compare them.

You have to tell me once, when you make each set.

Client Code

```
#include "set.h"

bool sameInt( void const *v1, void const *v2 )
{
    return *(int const *)v1 == *(int const *)v2;
}

int main()
{
    Set *s = makeSet( sizeof( int ), sameInt );
    insertValue( s, (int []){ 2 } );
    insertValue( s, (int []){ 3 } );
    insertValue( s, (int []){ 5 } );
    insertValue( s, (int []){ 7 } );
    insertValue( s, (int []){ 11 } );

    for ( int i = 0; i <= 12; i++ )
        if ( containsValue( s, &i ) )
            ...

    destroySet( s );
}
```

Here's how to compare two ints.

Make a set and tell it how to work with ints.

We have to pass values by address now (using array syntax to get the address of just one literal int).