

Security

CSC 230 : C and Software Tools
NC State Department of Computer
Science

Topics for Today

- Bad programs
 - Path manipulation
 - Log forging
 - Buffer overflow, exploited
- Encryption
- Sources of randomness

Island of Bad Programs

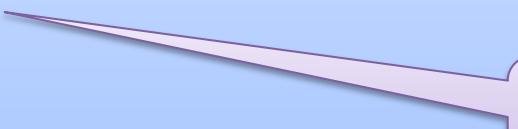
- C gives us a lot of power to write high-performance code
 - It's like programming without training wheels
- But, with great power ...
- It's easy to write programs with leaks, buffer overflows and security vulnerabilities
- I remember something about what causes this:
 - Programming errors
 - Failure to validate inputs
 - Inadequate protection for secret information
 - Poor understanding of the operating environment

Island of Bad Programs

```
/** Read an entire block of data from fp.  
 Return pointer to the block on success,  
 or NULL on failure */  
char *getBlock( FILE *fp ) {  
    char *buf = (char *)malloc( BLOCK_SZ );  
    if (!buf)  
        return NULL;  
    if (fread( buf, 1, BLOCK_SZ, fp ) != BLOCK_SZ )  
        return NULL;  
    else  
        return buf;  
}
```

Island of Bad Programs

```
/** Read an entire block of data from fp.  
   Return pointer to the block on success,  
   or NULL on failure */  
char *getBlock( FILE *fp ) {  
    char *buf = (char *)malloc( BLOCK_SZ );  
    if (!buf)  
        return NULL;  
    if (fread( buf, 1, BLOCK_SZ, fp ) != BLOCK_SZ )  
        return NULL;  
    else  
        return buf;  
}
```



I think you might have a
memory leak here.

- Be sure to clean up resources on all paths out of a function

Island of Bad Programs

```
/** Read the contents of the named text file into
   the given buffer.  If successful, return true.*/
bool getFile( char *name, int cap, char *buffer ) {
    FILE *fp = fopen( name, "r" );
    if ( !fp )
        return false;
    int len = fread( buffer, 1, cap, fp );
    return len > 0;
}
```

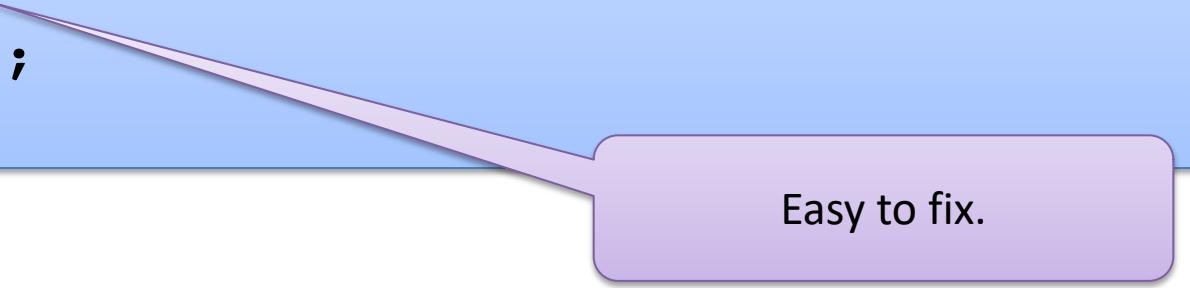
Island of Bad Programs

```
/** Read the contents of the named text file into  
the given buffer. If successful, return true.*/  
bool getFile( char *name, int cap, char *buffer ) {  
FILE *fp = fopen( name, "r" );  
if ( !fp )  
    return false;  
int len = fread( buffer, 1, cap, fp );  
return len > 0;  
}
```

You got yourself a file leak
right here.

Island of Bad Programs

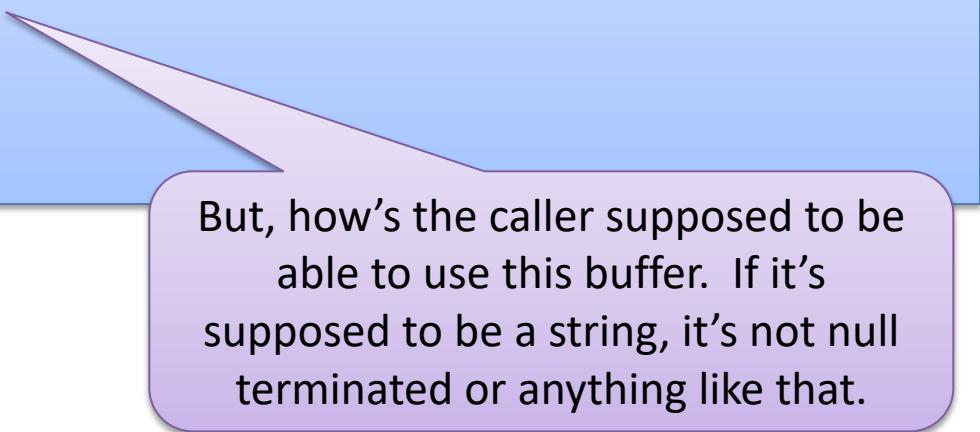
```
/** Read the contents of the named text file into  
the given buffer. If successful, return true.*/  
bool getFile( char *name, int cap, char *buffer ) {  
FILE *fp = fopen( name, "r" );  
if ( !fp )  
    return false;  
int len = fread( buffer, 1, cap, fp );  
fclose( fp );  
return len > 0;  
}
```



Easy to fix.

Island of Bad Programs

```
/** Read the contents of the named text file into  
the given buffer. If successful, return true.*/  
bool getFile( char *name, int cap, char *buffer ) {  
FILE *fp = fopen( name, "r" );  
if ( !fp )  
    return false;  
int len = fread( buffer, 1, cap, fp );  
fclose( fp );  
return len > 0;  
}
```



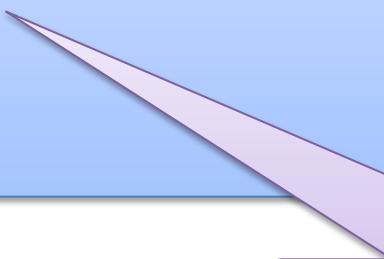
But, how's the caller supposed to be able to use this buffer. If it's supposed to be a string, it's not null terminated or anything like that.

Island of Bad Programs

```
/** Read a sequence of values from the user,
   returning the sequence and setting len on
   success, and returning NULL on failure. */
int *getValues( int *len ) {
    scanf( "%d", len );
    int *seq = (int *)malloc( *len * sizeof( int ) );
    for ( int i = 0; i < *len; i++ )
        if ( scanf( "%d", seq + i ) != 1 )
            return NULL;
    return seq;
}
```

Island of Bad Programs

```
/** Read a sequence of values from the user,  
    returning the sequence and setting len on  
    success, and returning NULL on failure. */  
int *getValues( int *len ) {  
    scanf( "%d", len );  
    int *seq = (int *)malloc( *len * sizeof( int ) );  
    for ( int i = 0; i < *len; i++ )  
        if ( scanf( "%d", seq + i ) != 1 )  
            return NULL;  
    return seq;  
}
```



There's a potential memory leak.

Island of Bad Programs

```
/** Read a sequence of values from the user,
   returning the sequence and setting len on
   success, and returning NULL on failure. /
int *getValues( int *len ) {
    scanf( "%d", len );
    int *seq = (int *)malloc( *len * sizeof( int ) );
    for ( int i = 0; i < *len; i++ )
        if ( scanf( "%d", seq + i ) != 1 ) {
            free( seq );
            return NULL;
        }
    return seq;
}
```

OK. Fixed.

Island of Bad Programs

```
/** Read a sequence of values from the user,
   returning the sequence and setting len on
   success, and returning NULL on failure. /
int *getValues( int *len ) {
    scanf( "%d", len );
    int *seq = (int *)malloc( *len * sizeof( int ) );
    for ( int i = 0; i < *len; i++ )
        if ( scanf( "%d", seq + i ) != 1 ) {
            free( seq );
            return NULL;
        }
    return seq;
}
```

What if this fails?

Island of Bad Programs

```
/** Read a sequence of values from the user,  
    returning the sequence and setting len on  
    success, and returning NULL on failure. /  
int *getValues( int *len ) {  
    if ( scanf( "%d", len ) != 1 )  
        return NULL;  
    int *seq = (int *)malloc( *len * sizeof( int ) );  
    for ( int i = 0; i < *len; i++ )  
        if ( scanf( "%d", seq + i ) != 1 )  
            free( seq );  
    return seq;  
}
```

Fixed

Still something to think about.
What if the user enters a
really, really big number?

Island of Bad Programs

```
/** Read a sequence of values from the user,
   returning the sequence and setting len on
   success, and returning NULL on failure. /
int *getValues( int *len ) {
    if ( scanf( "%d", len ) != 1 )
        return NULL;
    int *seq = (int *)malloc( *len * sizeof( int ) );
    for ( int i = 0; i < *len; i++ )
        if ( scanf( "%d", seq + i ) != 1 ) {
            free( seq );
            return NULL;
        }
    return seq;
}
```



I'd call all of these
programming
errors.

Hurting Lots of People

- Should we care about this kind of error?
 - When a user crashes a program, aren't they just hurting themselves?
 - ... well ... no. Not necessarily.
- We will develop software that:
 - serves lots of users
 - runs with more privileges than its users
- Vulnerabilities can lead to:
 - Denial of service for lots of users
 - Exposure of secret information

More Tools for Worse Programs

- There's a powerful, dangerous function

```
int system( const char *command );
```

- Give it a command, it will run it via the shell.
- This is a convenient, dangerous way to do lots of things you don't know how to do directly in C.

```
// I don't remember how to remove a file
system( "rm temp.txt" );
```

```
// I'm too lazy to write code to print out a file
system( "cat output.txt" );
```

More Tools for Worse Programs

- A function like `system()` is easy to use ... and easy to abuse.
- How about this code to print out a user-selected file?

```
printf( "File to view: " );
fgets( filename, sizeof(filename), stdin );

char cmd[ 1024 ] = "cat ";
strncat( cmd, filename, sizeof( cmd ) );

system( cmd );
```

More Tools for Worse Programs

- A function like `system()` is easy to use ... and easy to abuse.
- How about this code to print out a user-selected file?

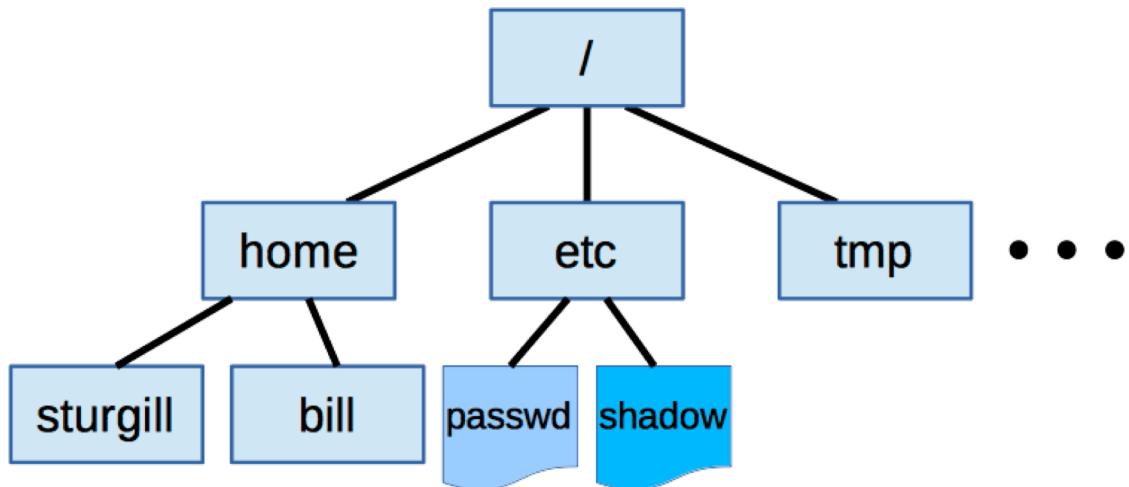
```
printf( "File to view: " );
fgets( filename, sizeof(filename), stdin );

char cmd[ 1024 ] = "cat ";
strncat( cmd, filename, sizeof( cmd ) );

system( cmd );
```

That's an error right there,
but it's not the worst one.

More Tools for Worse Programs



```
printf( "File to view: " );
fgets( filename, sizeof(filename), stdin );

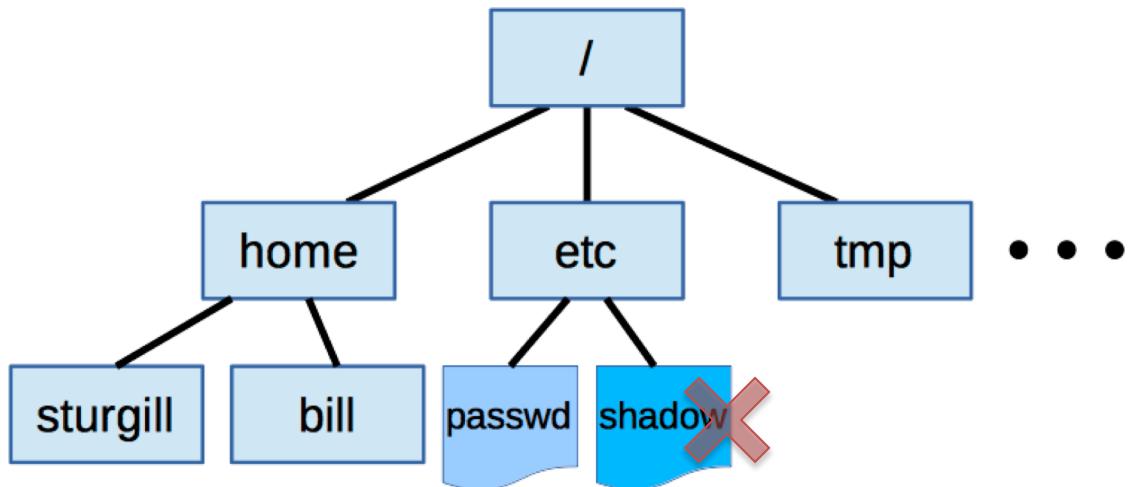
char cmd[ 1024 ] = "cat ";
strncat( cmd, filename, sizeof(cmd)-strlen(cmd)-1 );

system( cmd );
```

If the user types
/etc/shadow

cat /etc/shadow

More Tools for Worse Programs



These are examples of
Path Manipulation

```
printf( "File to view: " );
fgets( filename, sizeof(filename), stdin );
```

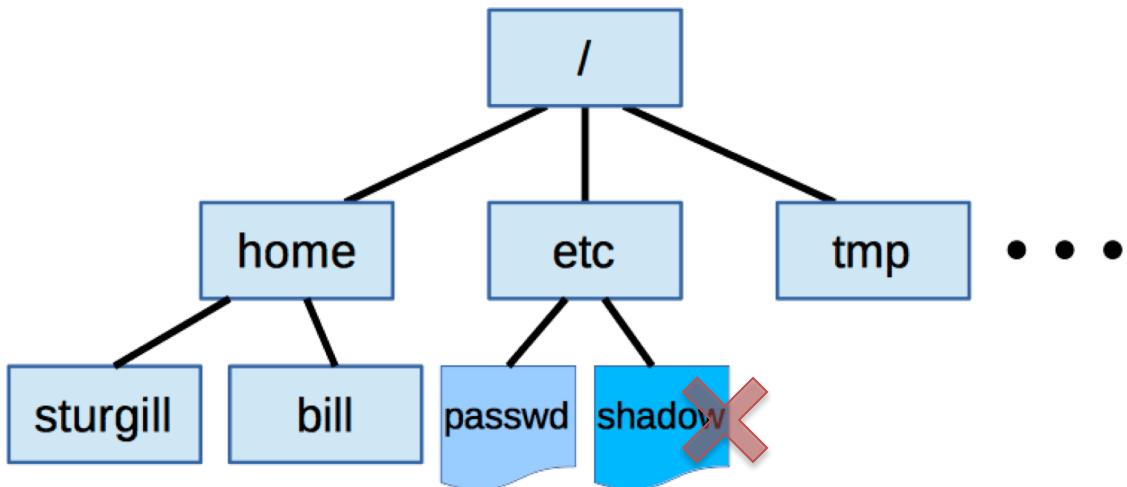
```
char cmd[ 1024 ] = "cat ";
strncat( cmd, filename, sizeof( cmd ) );
```

```
system( cmd );
```

Or, worse, what if they type
a.txt; rm /etc/shadow

cat a.txt; rm /etc/shadow

More Tools for Worse Programs



```
printf( "File to view: " );
fgets( filename, sizeof(filename), stdin );

char cmd[ 1024 ] = "cat ";
strncat( cmd, filename, sizeof( cmd ) - strlen( cmd ) - 1 );

system( cmd );
```

I'd call this failure
to validate inputs.

Validating Inputs

- Validate all inputs; don't rely on clients having done so
- Identify special (meta) characters and escape them consistently during input validation
- Use well-established, library functions to check for (a) legal URLs (b) legal filenames/pathnames (c) legal UTF-8 strings, ...
- Use white listing (rather than black listing)
- Authenticate that the communication is from correct person

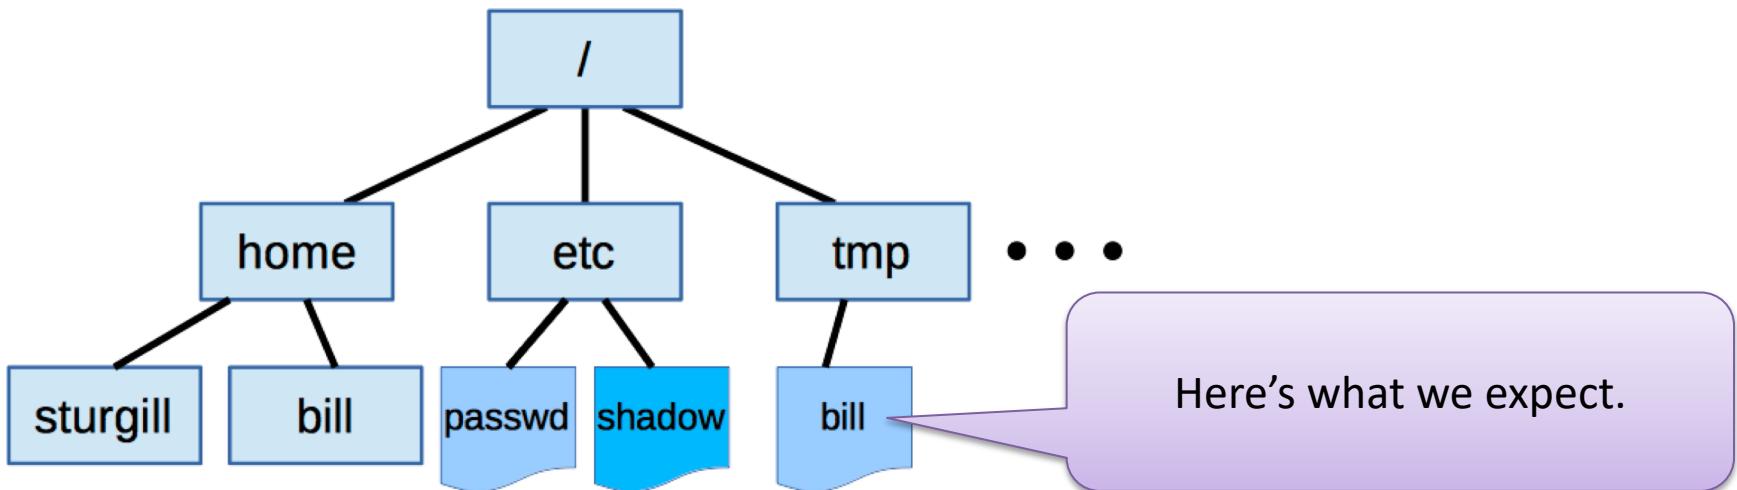
More Tools for Worse Programs

- A standard library function, `getenv()` will give us the value of environment variables.
- Useful, but not to be trusted.

```
...
// Make a temporary file for this user.
char *user = getenv( "USER" );
char path[ 1024 ];
snprintf( path, sizeof(path), "/tmp/%s", user );

FILE *fp = fopen( path, "w" );
...
```

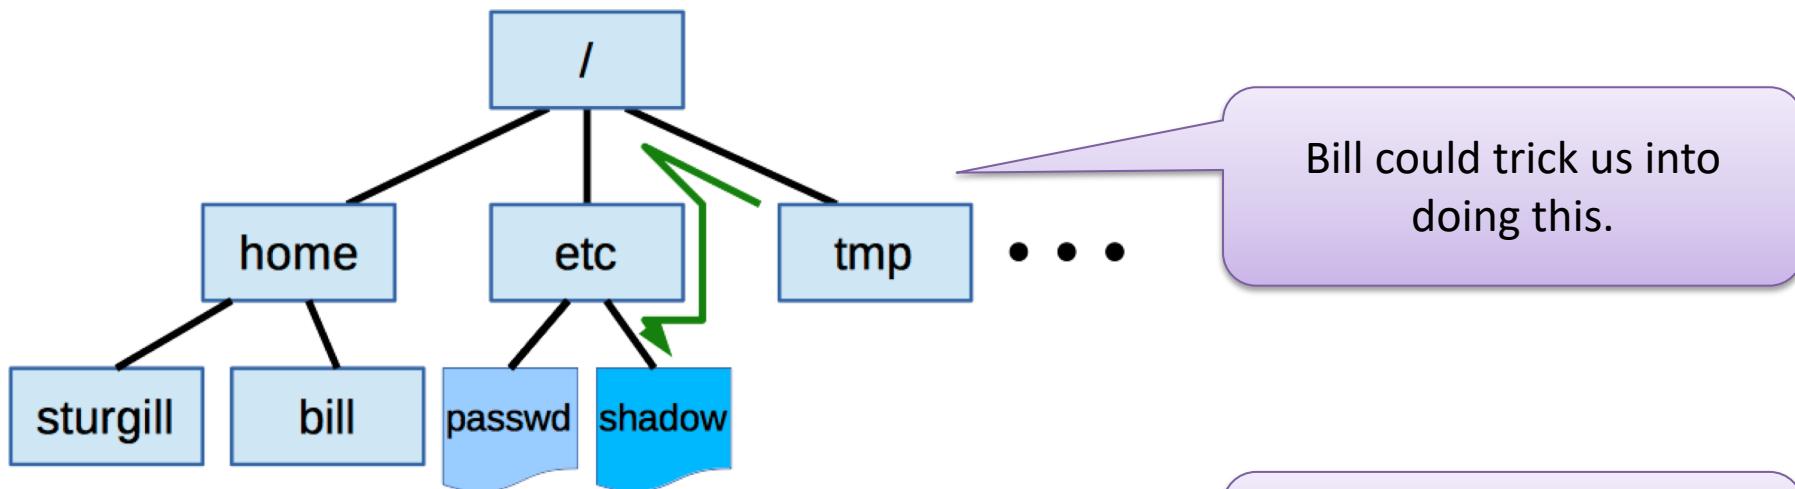
More Tools for Worse Programs



```
...
// Make a temporary file for this user.
char *user = getenv( "USER" );
char path[ 1024 ];
snprintf( path, sizeof(path), "/tmp/%s", user );

FILE *fp = fopen( path, "w" );
...
```

More Fun with Path Manipulation



```
...
// Make a temporary file for this user
char *user = getenv( "USER" );
char path[ 1024 ];
snprintf( path, sizeof(path), "/tmp/%s", user );

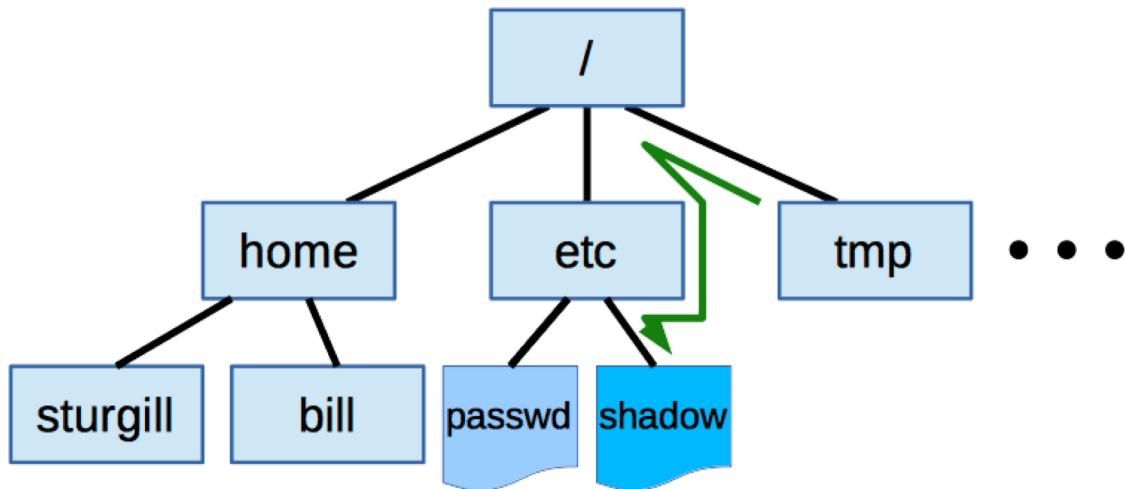
FILE *fp = fopen( path, "w" );
...

```

bill could do this:
\$ USER=..../etc/shadow

Overwrite
/tmp/..../etc/shadow
Ouch!

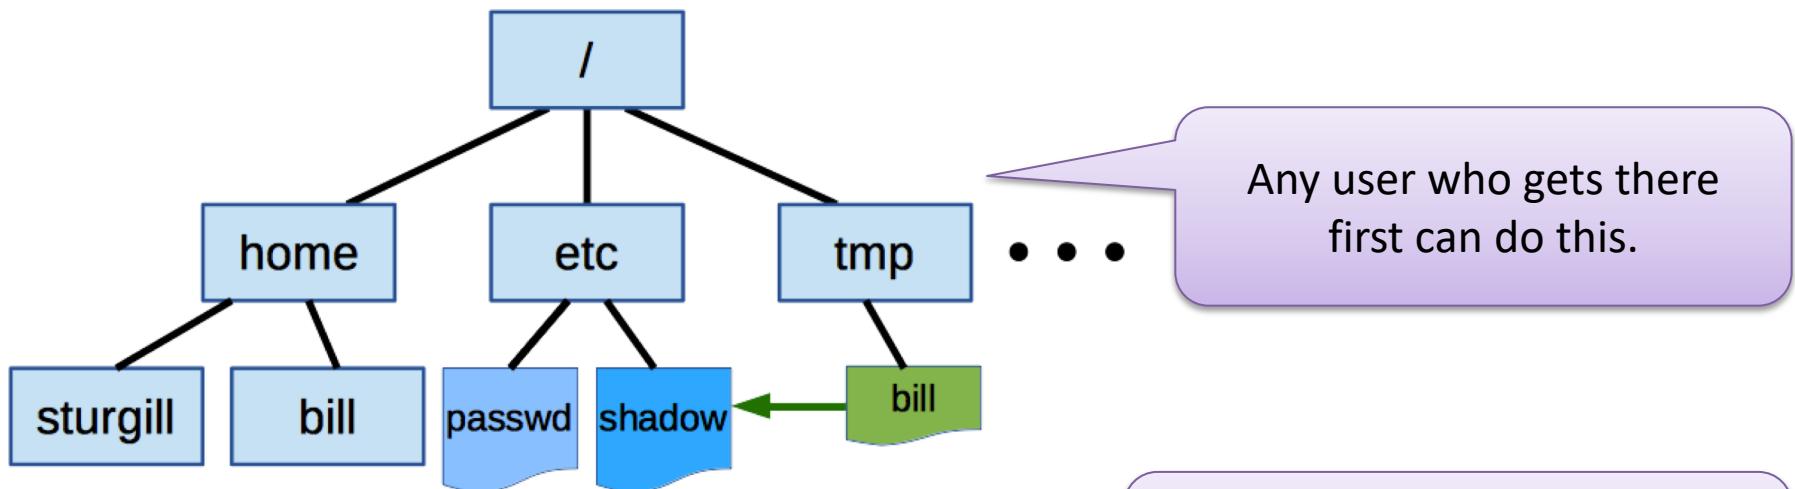
More Fun with Path Manipulation



```
...
// Make a temporary file for this user.
char *user = getenv( "USER" );
char path[ 1024 ];
snprintf( path, sizeof(path),
FILE *fp = fopen( path, "w" );
...
```

I'd call this **Inadequate Understanding of Operating Environment.**

More Fun with Path Manipulation



Any user who gets there first can do this.

```
...  
// Make a temporary file for this user  
char *user = getenv( "USER" );  
char path[ 1024 ];  
snprintf( path, sizeof(path), "/tmp/%s", user );  
  
FILE *fp = fopen( path, "w" );  
...
```

Any user can:
\$ ln -s /etc/shadow /tmp/bill

Now, this will overwrite
/etc/shadow
via the symbolic link.

Logging

- Sensitive applications can use logs to record:
 - Startup and configuration actions
 - Important events
 - Errors that occur
- Structured logs can help us detect and diagnose problems afterwards
- But, manipulating logs offers a way to “sow confusion”

Inviting Log Forging

```
char str[ 1000 ];
int len;
printf( "cmd> " );
if ( fgets( str, sizeof(str), fp ) ) {
    if ( valid( str ) ) {
        ...
    } else
        fprintf( log, "LOG: invalid string = %s\n", str );
}
...
...
```

- Great. We can tell when something went wrong.
- ... and maybe other nonsense we didn't expect.

Log Forging

- If the a user enters:

```
cmd> complete-nonsense
```

- The log does its job:

```
LOG: invalid string = complete-nonsense
```

- But, an attacker may enter misleading strings:

```
cmd> complete-nonsense  
authenticated = Darth Vader
```

```
LOG: User
```

- This is going to create some misleading log entries:

```
LOG: invalid string = complete-nonsense  
LOG: user authenticated = Darth Vader
```

Log Forging

- If the user enters:

```
cmd> complete-nonsense
```

- The log does its job:

```
LOG: invalid string = complete-nonsense
```

A bunch of spaces here.

- But, an attacker may enter misleading strings:

```
cmd> complete-nonsense  
authenticated = Darth Vader
```

- This is going to create some misleading entries:

```
LOG: invalid string = complete-nonsense  
LOG: user authenticated = Darth Vader
```

LOG: User

Makes this look like a line break, if you're not careful.

Buffer Overflow ... Revisited

- Buffer overflow will crash your program
 - If you're lucky.
- If not, it's something an attacker can use to get your system
 - To make it fail in exactly the way she wants.

```
void f() {  
    int a[10];  
    a[20] = 3;  
}
```

Example: Buffer Problem

```
bool passwd_ok = false;
char passwd[ 10 ];

printf( "Password: " );
scanf( "%s", passwd );

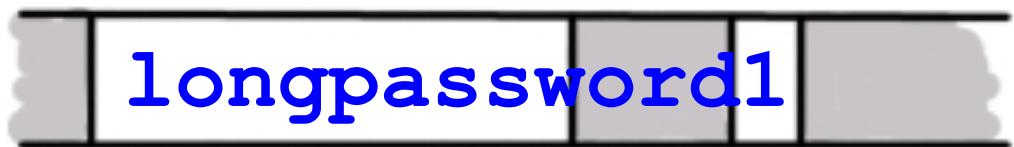
if ( strcmp( passwd, "joshua" ) == 0 )
    passwd_ok = true;

if ( passwd_ok ) {
    printf( "Welcome, the system is yours.\n" );
}
```

- Layout in memory:

passwd

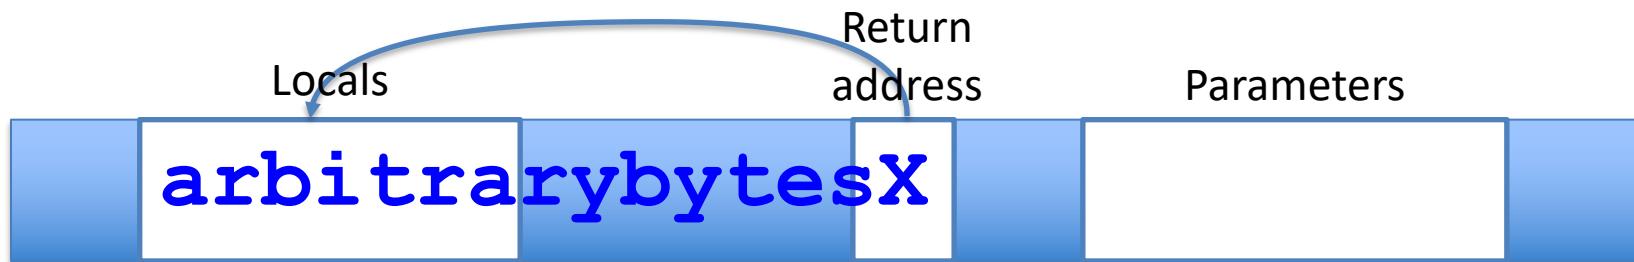
passwd_ok



- **passwd** buffer overflowed, overwriting **passwd_ok** flag
 - Any sufficiently long password will be accepted!

Stack Attacks

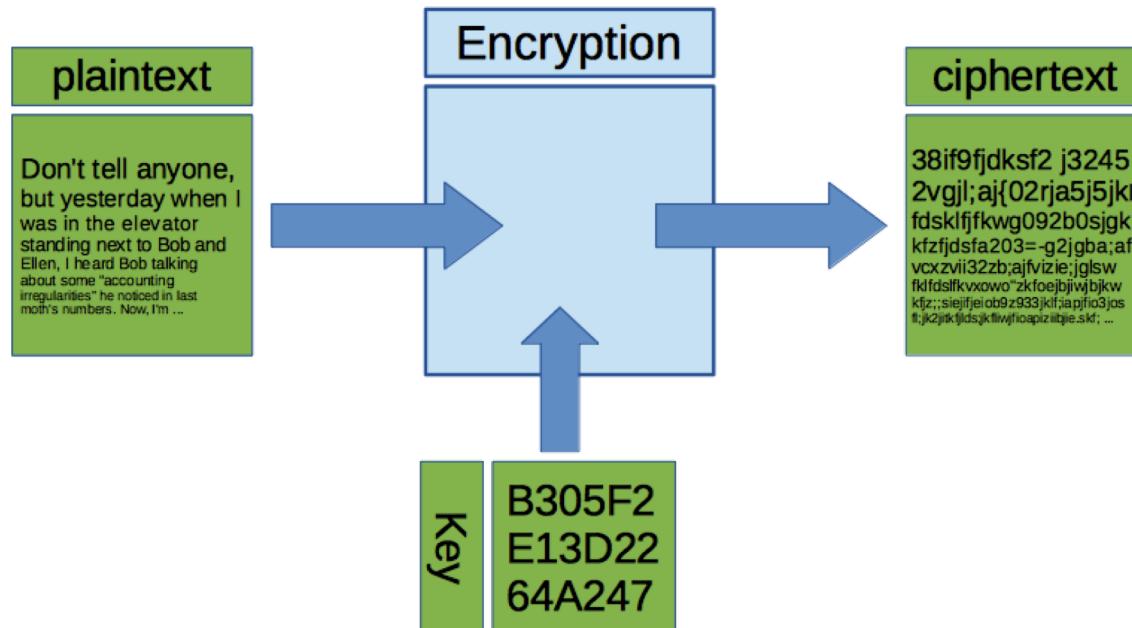
- When a function is called...
 - parameters are pushed on stack
 - return address pushed on stack
 - called function puts local variables on the stack
- Memory layout



- Problems?
 - Return to address X which may execute arbitrary code
 - Maybe even code they just wrote into memory.

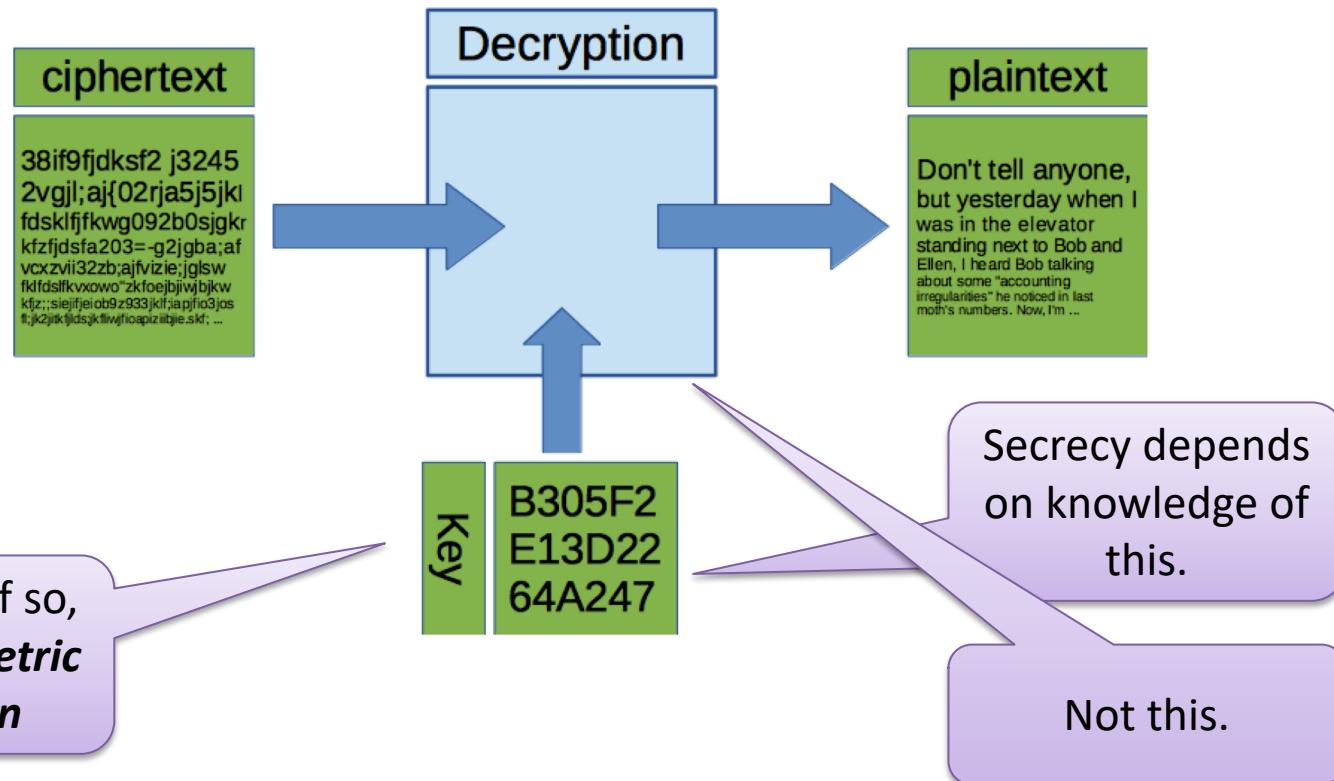
Encryption

- *Encryption* : Turning *plaintext* into something that looks like random garbage (*ciphertext*), using a *key*

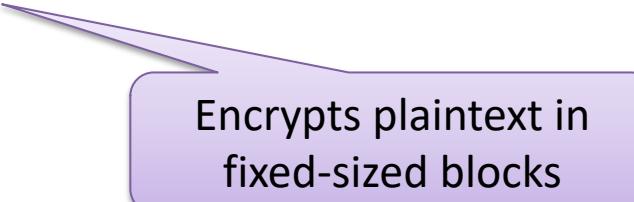


Decryption

- Decryption: turning ciphertext back into plaintext, based on a key.



Symmetric Encryption Techniques

- There are a variety of symmetric encryption algorithms
 - DES : 56-bit key *block cipher*. Good enough ... for 1977.
 - Many newer algorithms supporting a larger *keyspace*
 - 3DES, AES
- 
- Encrypts plaintext in fixed-sized blocks

Fun with Encryption

- The openssl command knows a lot of encryption techniques

```
$ openssl list-cipher-algorithms
AES-128-CBC
AES-128-CBC-HMAC-SHA1
... Big list of encryption techniques.
DES => DES-CBC
DES-CBC
... More techniques
AES-128-CBC
... Even more
```

- It's a command-line utility that implements the Secure Sockets Layer and related protocols

Fun with Encryption

- Let's give it a try:
 - Make a file message.txt containing anything you want.
 - Encrypt the file with openssl

```
$ openssl enc -DES -in message.txt -out message.bin
```

Do encryption / decryption.

Use DES.

Input and output files.

Fun with Encryption

- We get a new, encrypted file, message.bin
 - The original file is still there.
 - Are you brave enough to delete it? I am.

```
$ rm message.txt
```

- I hope I can get it back. Of course, that's what encryption is all about.

```
$ openssl enc -DES -d -in message.bin -out message.txt
```

- I have to provide the same key for this to work.

The -d is for decrypt.

A Bad Idea

- I don't like typing in my password twice.
 - Is there a way to automate this?
 - Yes. But some ways are bad ideas.

Don't do this.

```
$ openssl enc -DES -pass pass:abc123 -in message.txt -out message.bin
```

- Anyone on the system can potentially see this password
- This is really about **knowing your execution environment.**

Encryption from C

- We could run a command like this using `system()`

```
system( "openssl enc -DES -d -pass pass:abc123 ..." );
```

- But, that might expose our password
 - ... and it could be slow.

Encryption from C

- There are libraries to encrypt / decrypt for you.
- libcrypt is available on many Unix/Linux systems (although it's not part of the C standard)
- You have to include the **rpc/des_crypt.h** header.
- This gives you a few functions:

8-byte key.

Data to encrypt/decrypt, some multiple of 8 bytes in length.

```
ecb_crypt( char *key, char *data,  
           unsigned int len, unsigned int mode );
```

Length of data to encrypt

DES_ENCRYPT or
DES_DECRYPT

Key Preparation

- The key has a parity bit for each byte.
 - Before you can use it, you need to set up its parity bits.
 - There's a function that will do this for you.

```
des_setparity( key );
```

8-byte key

Encrypt / Decrypt Example

```
#include <rpc/des_crypt.h>
```

```
char key[] = "abcd1234";  
des_setparity( key );
```

Prepare the 8-byte key.

```
char message[] = "Here, an example message";  
int msgLen = strlen( message );
```

Make a 24-byte message.

```
ecb_crypt( key, message, msgLen, DES_ENCRYPT );
```

Encrypt the message (in place)

```
ecb_crypt( key, message, msgLen, DES_DECRYPT );
```

Decrypt the message (in place)

Linking with libcrypt

- These functions are in a library, libcrypt
- To build an executable, you need to link with this library.

```
$ gcc -Wall -std=c99 -g des.c -o des -lcrypt
```

Link with libcrypt

Actually, my executable
compiled OK without this.

Software Security

- Think about security up-front
- Consider security as functionality rather than hidden part of system
- Design and test with security in mind
- Protect your secrets and paths of communication
 - Cryptography
- Program defensively
 - Validate inputs
 - Check buffers and bounds
- Verification and Validation
 - Test! Think maliciously! How could you attack a system?
 - Use tools that support identifying security vulnerabilities.

Hiding Information

- It can be hard to hide information in a program:

```
char passwd[ 64 ];  
printf( "Enter password: " );  
scanf( "%63s", passwd );  
if ( strcmp( passwd, "Albuquerque" ) == 0 ) {  
    ...  
}
```

- Not very secure if this is open source.
- But, even if it's not ...

```
$ strings program  
Enter password:  
%63s  
Albuquerque  
...
```

Hiding Information

- It can be hard to hide information in a program:

```
char passwd[ 64 ];
printf( "Enter password: " );
scanf( "%63s", passwd );
if ( strcmp( passwd, "Albuquerque" ) == 0 ) {
    ...
}
```

- Not very secure if this is open source.
- But, even if it's not ...

```
$ strings program
Enter password:
%63s
Albuquerque
...
```



I'd call this Inadequate Protection of Secret Information.

Depending on Randomness

- We don't usually depend on hiding secrets in an executable
 - Usually, we want to choose our own secrets
 - ... and store them as part of our configuration
- This can benefit from a source of randomness
 - How about this:

```
char passwd[ 15 ];
for ( int i = 0; i < sizeof( passwd ) - 1; i++ )
    passwd[ i ] = 'a' + rand() % 26;
passwd[ sizeof( passwd ) - 1 ] = '\0';
```

Depending on Randomness

- Bad idea. This is a pseudo-random sequence.
It's the same from execution to execution.
- How about this instead?

```
strand( time( NULL ) );
char passwd[ 15 ];
for ( int i = 0; i < sizeof( passwd ); i++ )
    passwd[ i ] = 'a' + rand() % 26;
passwd[ sizeof( passwd ) - 1 ] = '\0';
```

- Not much better. The space of possible times
is still quite small.

More Random Randomness

- Unix-like systems offer `/dev/random`
 - A sequence of random bytes
 - Based on pseudorandom processes and environmental noise
- Need 16 bytes of randomness? No problem.

```
$ dd count=16 bs=1 if=/dev/random of=randbits.bin
```

- Random enough for you?

```
$ hexdump -C randbits.bin
00000000  de 92 b9 5e 52 75 e3 0b  f3 50 29 20 83 3b ac 86
```