

# Performance and Optimization

CSC 230 : C and Software Tools

NC State Department of Computer  
Science

# Topics for Today

- Leftover Topic: Environment Variables
- Code coverage
- Optimization Techniques
  - Timing program execution
  - Profiling
  - Optimizing by Hand
  - Compiler support for optimization

# Environment Variables

- Your execution environment supports variables
  - We can access these from the shell

```
$ echo $HOME
/home/jerry
$ HOME=/home/linda
$ echo $HOME
/home/linda
```

- The shell uses these a lot for configuration and to customize behavior
  - HOME, PATH, SHELL, PS1, EDITOR, ...
- But, these aren't just available to the shell, all programs have environment variables.

# Using the Environment

- The standard library lets us get/set these variables
  - We need our good friend, `stdlib.h`
  - Then, we can get environment variable values

```
char *str = getenv( "HOME" );  
printf( "HOME = %s\n", str );
```

- And, we can change them

```
setenv( "HOME", "/some/other/place", 1 );
```

Variable to set.

new value.

Should I overwrite?

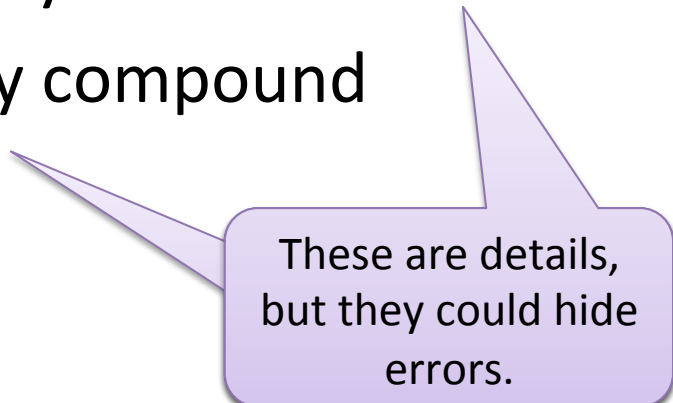
- These changes are local to the program, they go away when it exits.

# Testing Testing

- How do you know if your tests are adequate?
  - Well, there are different types of tests you can perform
    - Unit testing, integration testing, system testing, regression testing, acceptance testing, beta testing
    - black-box vs white box
  - Each type intended to expose a different cause or type of bug.
- For acceptance testing, we could make sure we have a test or two for every (functional) requirement
- For black box testing, we can test output on typical and boundary cases
  - To make sure we cover the range of possible inputs

# Code Coverage

- For white-box testing, we can make sure we cover the code itself.
- Ideally, all the source code should get a chance to run under test
  - Every statement executed
  - Every branch direction of every conditional
  - Every sub-expression of every compound conditional



These are details,  
but they could hide  
errors.

# Tools for Code Coverage

- The compiler will help us measure this
- Compile with the `-fprofile-arcs` and `-ftest-coverage` flags

```
gcc -Wall -std=c99 -fprofile-arcs \  
    -ftest-coverage bubble1.c -o bubble1
```

- Run your program all you like

```
./bubble1 < small_input  
./bubble1 < some_other_input
```

# Using gcov

- Notice, you get some extra files.

```
eos$ ls
bubble1      bubble1.c
bubble1.gcno bubble1.gcda
```

written during  
execution

written by the  
compiler

- Run **gcov** to extract the coverage information

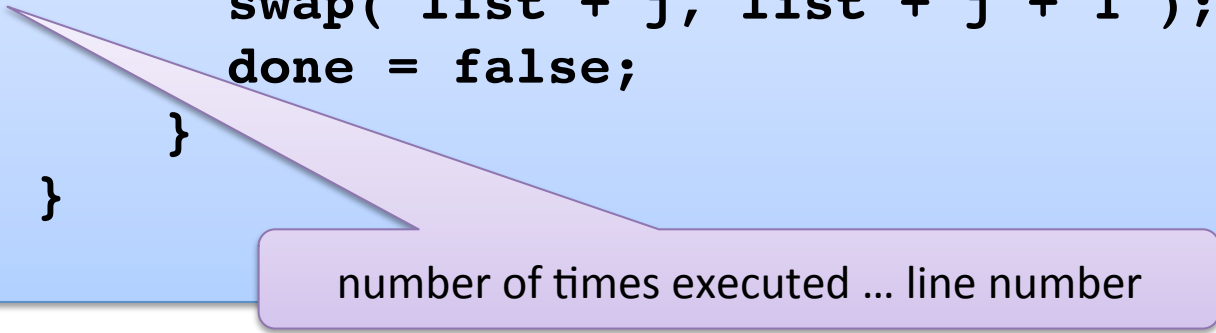
```
eos$ gcov bubble1.c
File 'bubble1.c'
Lines executed:66.67% of 36
bubble1.c:creating 'bubble1.c.gcov'
```



# Reading gcov Output

- In addition to the summary, gcov writes a text file giving coverage details.

```
2: 42: for ( int i=0; !done && i<len; i++ ) {
-: 43:     // If we make it through without ...
1: 44:     done = true;
-: 45:
5: 46:     for ( int j=0; j<len-i-1; j++ )
4: 47:         if ( list[j] > list[j+1] ) {
#####: 48:             swap( list + j, list + j + 1 );
#####: 49:             done = false;
-: 50:         }
-: 51:     }
1: 52: }
```



number of times executed ... line number

# Reading gcov Output

- gcov has some codes for reporting execution count

5	I was executed 5 times
-	I could never be executed (this doesn't count against coverage)
#####	I was never executed (but I could have been)

- Let's see if we can get code coverage up to 100 %

# Getting More from gcov

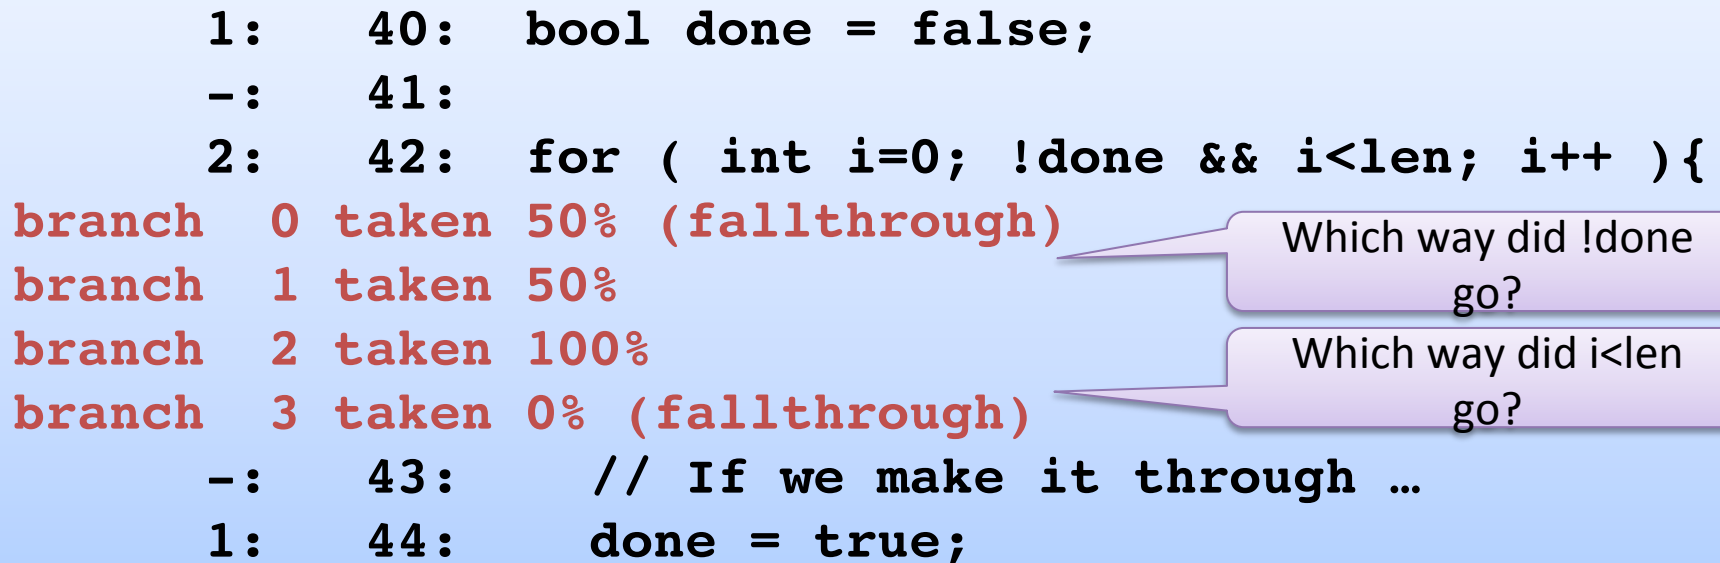
- By default, gcov tells you about statements executed
- It can also tell you how branches went

```
eos$ gcov -b bubble1.c
File 'bubble1.c'
Lines executed:66.67% of 36
Branches executed:100.00% of 18
Taken at least once:77.78% of 18
Calls executed:61.54% of 13
bubble1.c:creating 'bubble1.c.gcov'
```

# Reading gcov Output

- With the -b option, gcov reports branching directions for each part of each conditional.

```
1: 40: bool done = false;
-: 41:
2: 42: for ( int i=0; !done && i<len; i++ ){
branch 0 taken 50% (fallthrough)
branch 1 taken 50%
branch 2 taken 100%
branch 3 taken 0% (fallthrough)
-: 43: // If we make it through ...
1: 44: done = true;
```



The diagram illustrates the gcov output for a C++ for loop. The code is shown with line numbers and column numbers. The loop body contains a for loop with two conditions: `!done` and `i < len`. The gcov output shows four branches: branch 0 (taken 50%, fallthrough), branch 1 (taken 50%), branch 2 (taken 100%), and branch 3 (taken 0%, fallthrough). Two callouts are present: one pointing to branch 0 asking "Which way did !done go?" and another pointing to branch 3 asking "Which way did i < len go?".

# Performance and Optimization

- Performance depends on
  - Coding Style
  - Choice of programming language
  - Compiler and Options

**—Choice of Algorithm and  
Data Structure**

# Code Profiling

- We can measure our program's performance using the time command.

```
time ./program program-options
```

- This reports how long it took the program to execute, CPU time consumed, etc.
- It's fairly coarse grained, we'd like to know more about where our program is spending its time.
- Code profiling tools let us do that.

# Code Profiling

- **gprof** is a tool for profiling during execution
  - It counts the number of times each function is called
  - How much time we spend in each function
- To use gprof, we need to compile with the **-pg** option:

```
gcc -pg -Wall -std=c99 program.c -o program
```


# Generating Profile Information

- Then, you can run your program as normal.

```
./program program-options
```

- This also generates some profiling output

```
$ ls  
gmon.out      myProgram      myProgram.c
```



Where did this  
file come from?



# Examining Profile Output

- The gmon.out file is in a binary format.
- But, the gprof command will print it in a way that's human readable.

```
$ gprof ./myProgram  
<lots of output>
```

- Maybe it would be better to send this output to a file.

```
$ gprof ./myProgram > perf.txt
```

# Fun With Time and Gprof

- Let's run bubble4.c against the big\_input.txt input file.
  - This file is large enough to take a few seconds to bubble sort.
- First, let's time the execution
- Then, let's profile:
  - Compile with profiling support
  - Run to collect profiling results
  - Examine results with gprof

# DIY Optimizations : Code Motion

- You can optimize code yourself
  - By moving duplicate expression evaluation outside loops

```
for ( int i = 0; i < n; i++ )  
    for ( int j = 0; j < n; j++ )  
        a[ n * i + j ] = f() + g(i) + h(j);
```



```
int c1 = f();  
for ( int i = 0; i < n; i++ ){  
    int c2 = c1 + g(i);  
    int ri = n * i;  
    for ( int j = 0; j < n; j++ )  
        a[ ri + j ] = c2 + h(j);  
}
```

# DIY Optimizations : Strength Reduction

```
c1 = f();  
for ( int i = 0; i < n; i++ ){  
    c2 = c1 + g(i);  
    int ri = n * i;  
    for ( int j = 0; j < n; j++ )  
        a[ ri + j ] = c2 + h(j);  
}
```




```
c1 = f();  
*p = a;  
for ( int i = 0; i < n; i++ ){  
    c2 = c1 + g(i);  
    for ( int j = 0; j < n; j++ )  
        *p++ = c2 + h(j);  
}
```

# DIY Optimization : Expression Reuse

- Consider this neighborhood computation:

```
sum  = a[ (i-1)*m + j  ];  
sum += a[ (i+1)*m + j  ];  
sum += a[ i*m +      j-1 ];  
sum += a[ i*m +      j+1 ];
```



```
int ctr = i*m + j;  
sum  = a[ ctr - m ];  
sum += a[ ctr + m ];  
sum += a[ ctr - 1 ];  
sum += a[ ctr + 1 ];
```

# DIY Optimization : Inlining


```
double prod( int i, int j )  
{  
    double sum = 0.0;  
    for ( int k = 0; k<n; k++)  
        sum += a[i][k] * b[k][j];  
    return sum;  
}
```

```
for (i=0; i<n; i++)  
    for (j=0; j<n; j++)  
        c[i][j] = prod(i, j);
```

# DIY Optimization : Inlining

```
double prod( int i, int j )
{
    double sum = 0.0;
    for ( int k = 0; k<n; k++)
        sum += a[i][k] * b[k][j];
    return sum;
}
```

```
for (i=0; i<n; i++)
    for (j=0; j<n; j++) {
        double sum = 0.0;
        for (k=0; k<n; k++ )
            sum += a[i][k] * b[k][j];
        }
    c[i][j] = sum;
}
```



# Inlining Help from the Language

- C can inline your functions for you
  - Introduced in C99

```
inline double prod( int i, int j )  
{  
    double sum = 0.0;  
    for ( int k = 0; k<n; k++)  
        sum += a[i][k] * b[k][j];  
    return sum;  
}
```

```
for (i=0; i<n; i++)  
    for (j=0; j<n; j++)  
        c[i][j] = prod(i, j);
```



# Help from the Compiler

- Many of these refinements can be done automatically by the compiler
- It can try to optimize for execution speed
  - Options `-O`, `-O2` and `-O3`, in order of increasing optimization.
- Increasing optimization level can:
  - Increase code size
  - Possibly decreasing performance
- But, the compiler can't do everything

# DIY Optimization : Reordering Tests

```
if ( height > 84 )           // rare
    reallyTall();
else if ( height > 72 )      // unusual
    tall();
else                          // typical
    average();
```

- Place more frequent branches earlier


```
if ( height <= 72 )          // typical
    average();
else if ( height <= 84 )     // unusual
    tall();
else                          // rare
    reallyTall();
```

# DIY Optimization : Pass by Reference

- C lets us use value semantics with most types
  - ... maybe it's too easy.
  - Passing large objects by value can be expensive.

```
int f( struct BigRecord rec )  
{  
    int sum = rec.x + rec.y;  
    return sum;  
}
```

```
struct BigRecord a;  
...  
v = f( a );
```



# DIY Optimization : Pass by Reference

- Pass by reference/address can reduce overhead copying large objects
- Const lets the function say what it plans to do with the object.

```
int f( const struct BigRecord *rec )  
{  
    int sum = rec->x + rec->y;  
    return sum;  
}
```

```
struct BigRecord a;  
...  
v = f( &a );
```

